

La Serena School for Data Science

Manifold learning
and dimensionality
reduction

2024

Federica Bianco

University of Delaware

Rubin Observatory

Special thanks to A. Bayo

access this presentation at
https://slides.com/federicabianco/lssdss24_AEC

1/6

dimensionality reduction - why

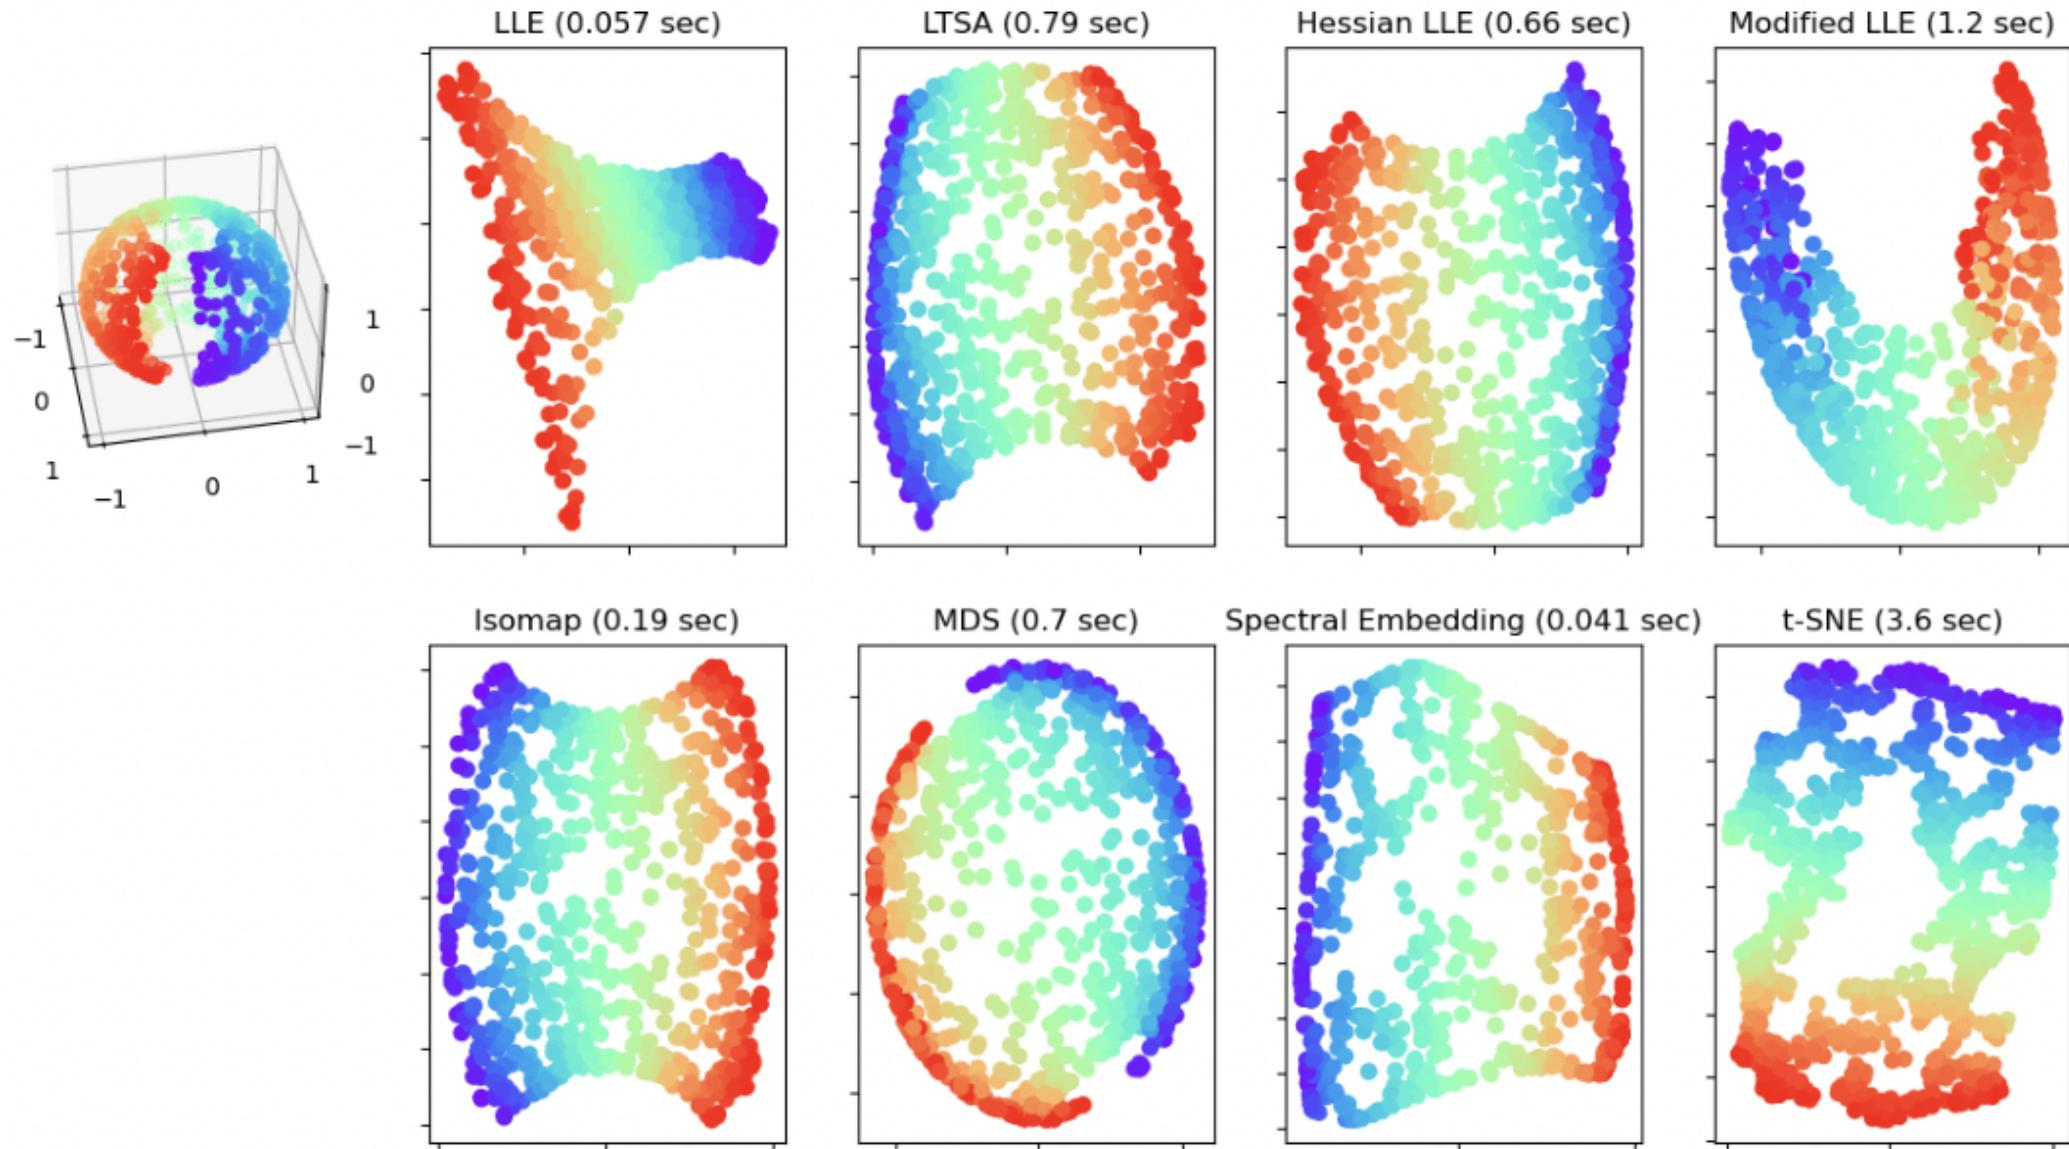
- Visualization
- Computational cost
- Course of dimensionality



2/6

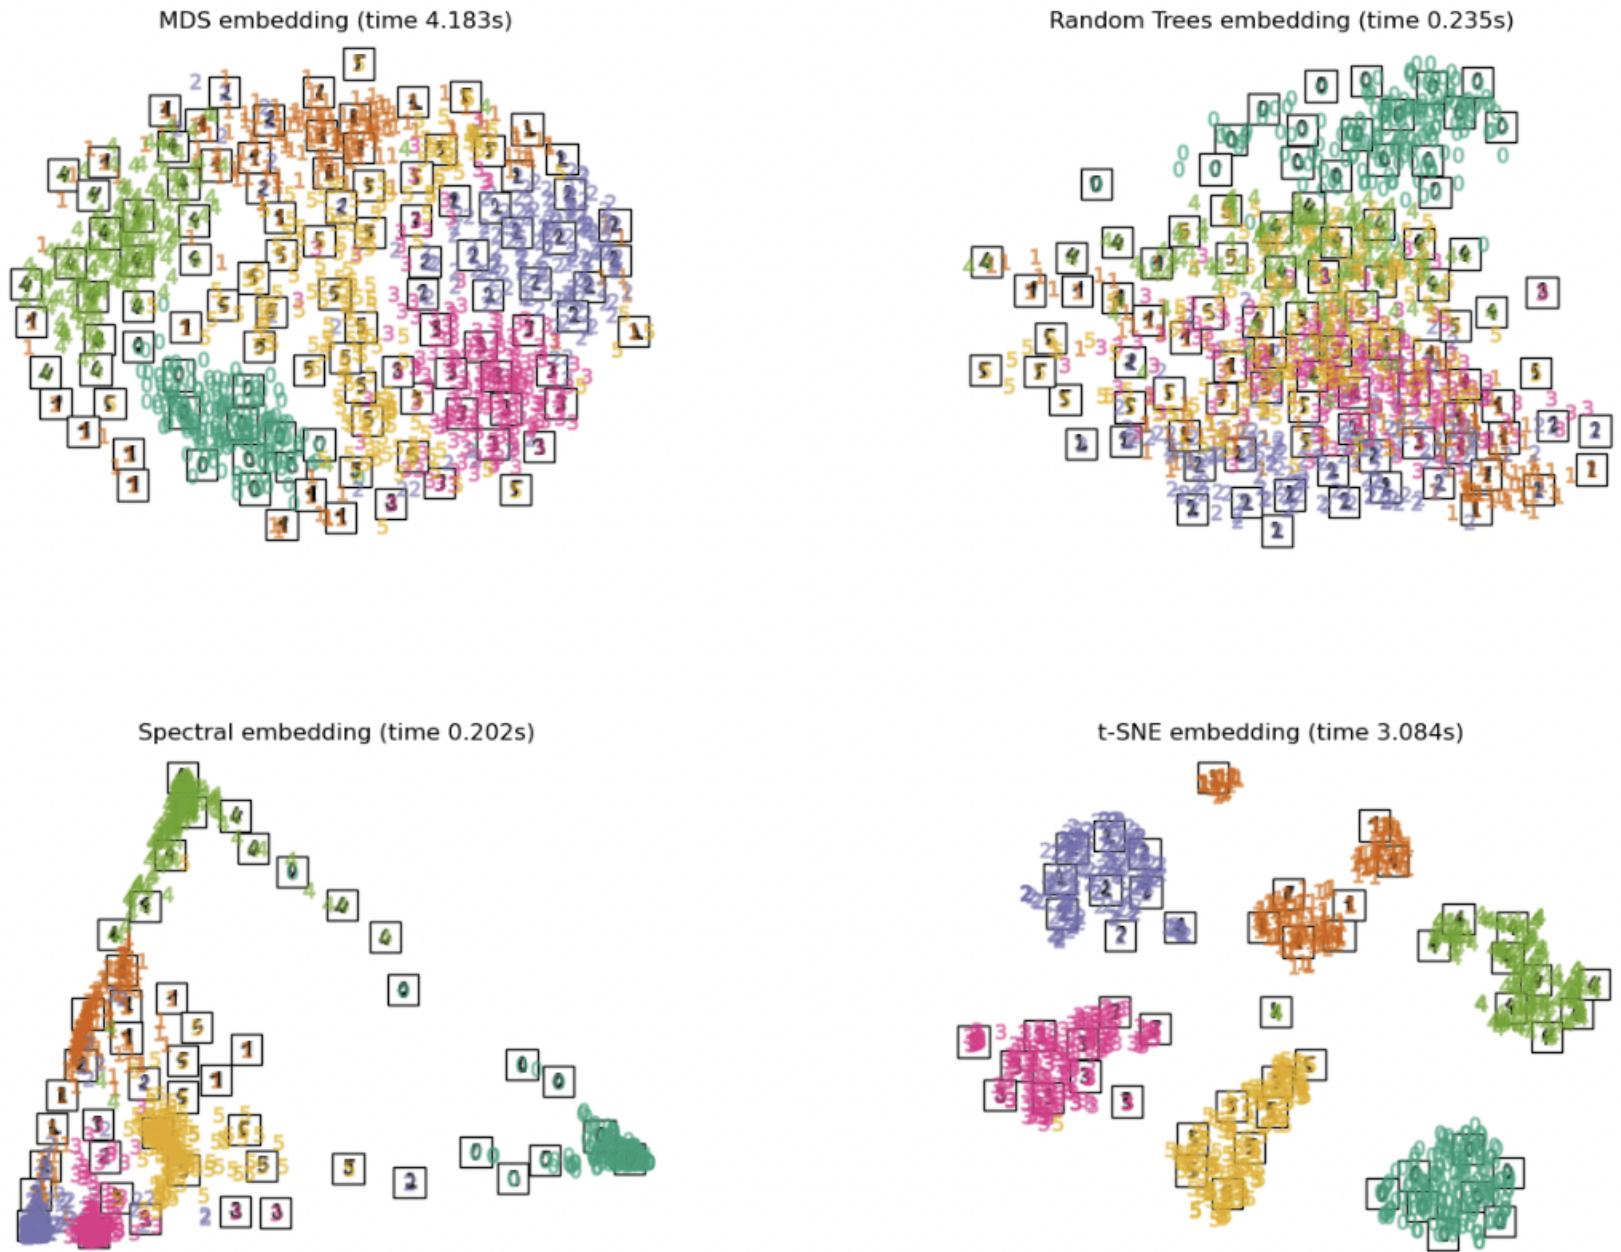
manifold lerning

Manifold Learning with 1000 points, 10 neighbors



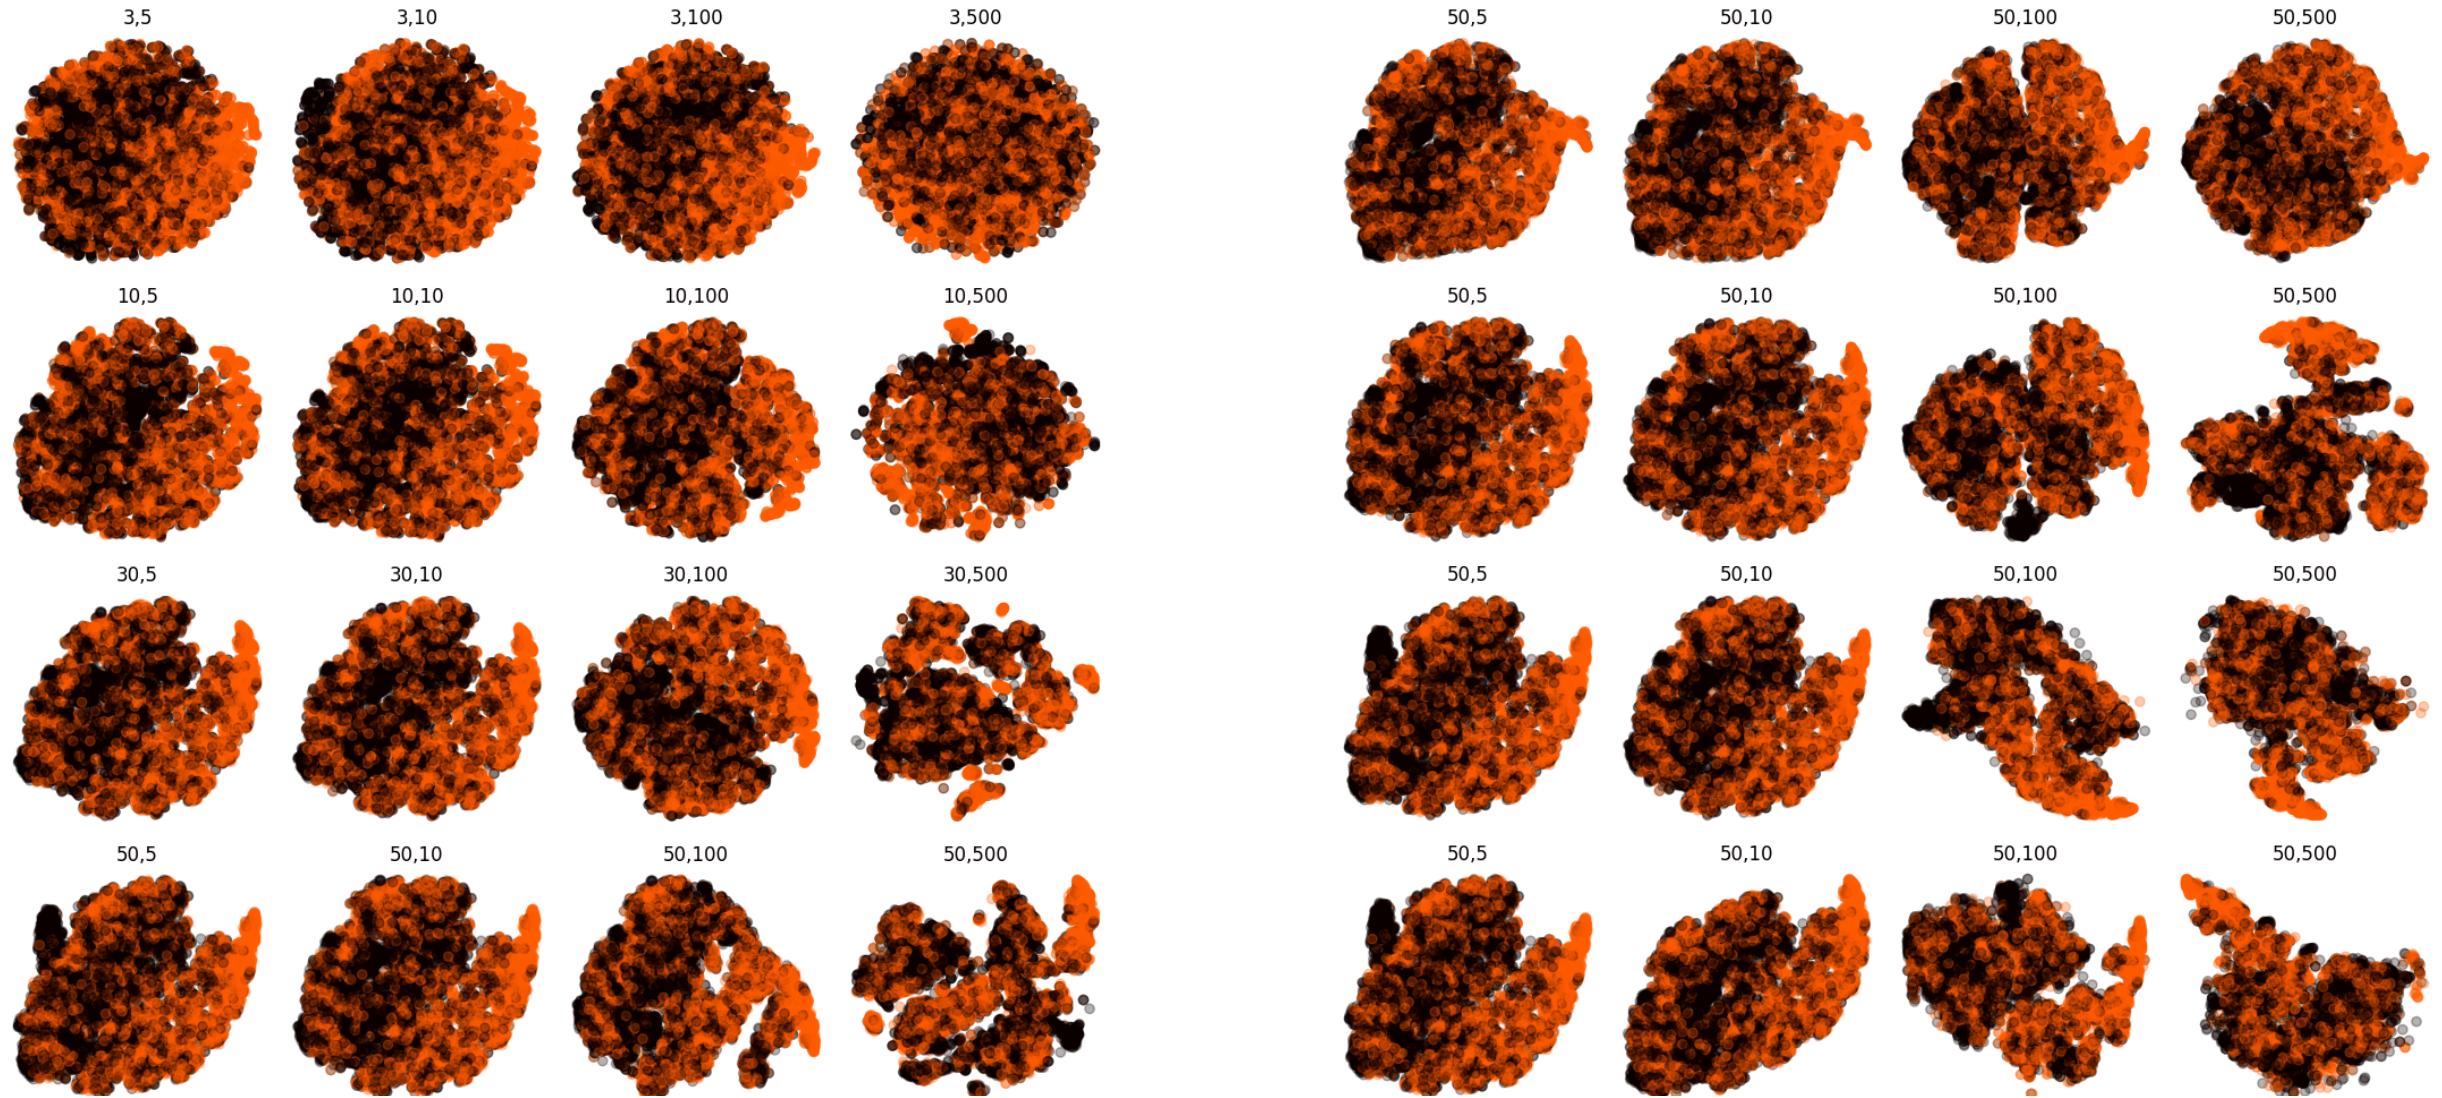
MANIFOLD LEARNING ON MNIST

https://scikit-learn.org/stable/auto_examples/manifold/plot_lle_digits.html



Higgs boson Kaggle dataset Training set of 250000 events, with an ID column, 30 feature columns, a weight column and a label column.

perplexity, early exaggeration

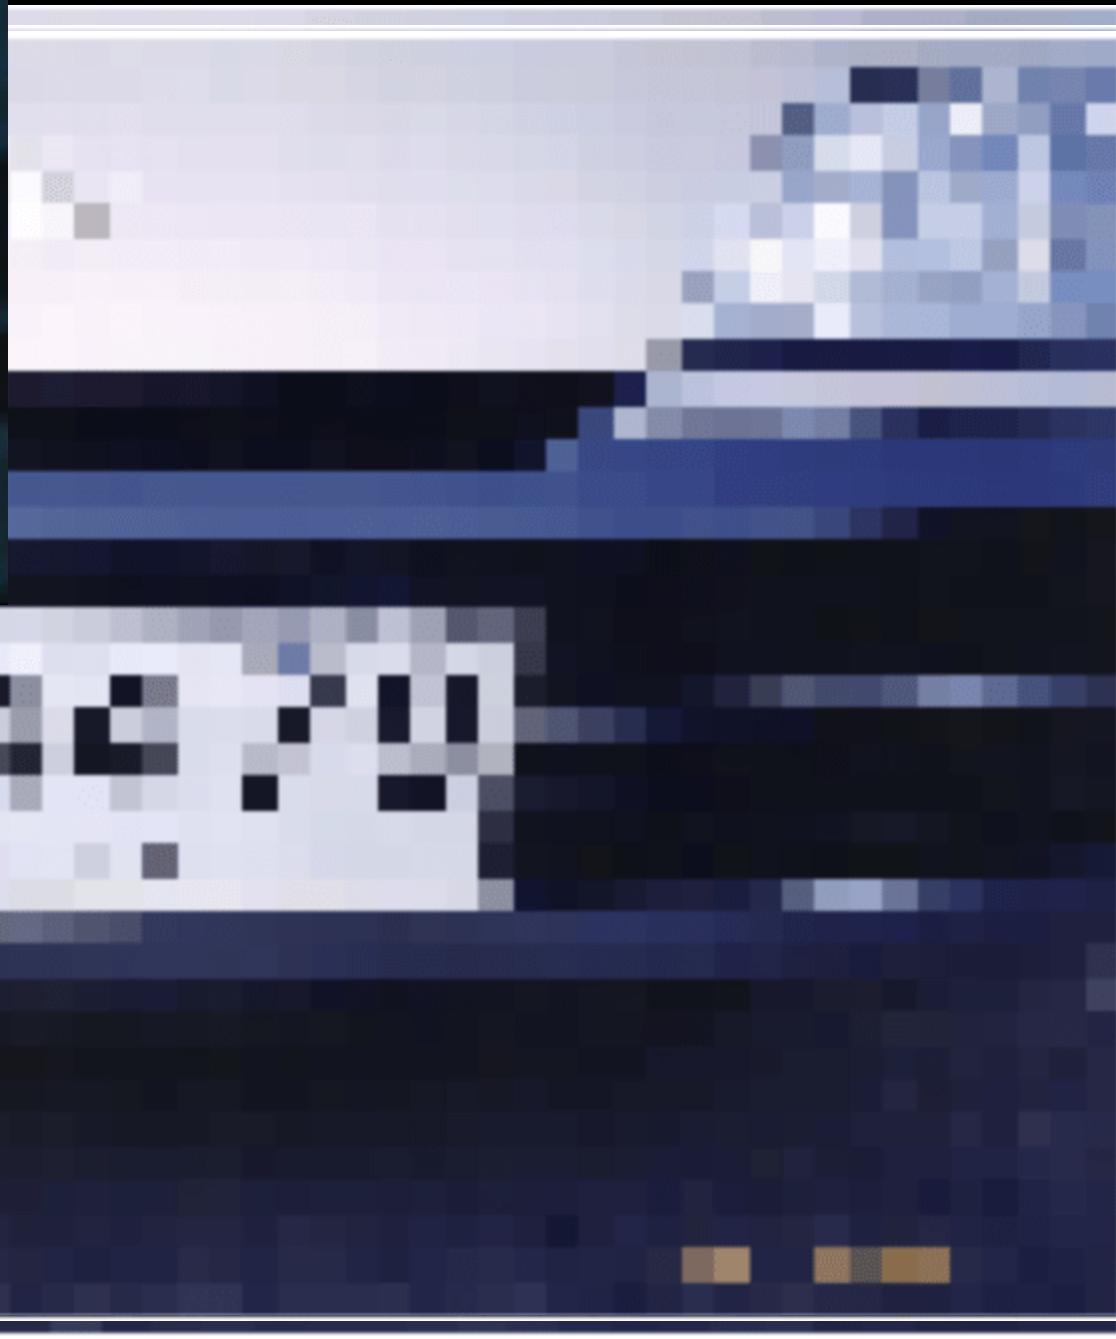




3/6

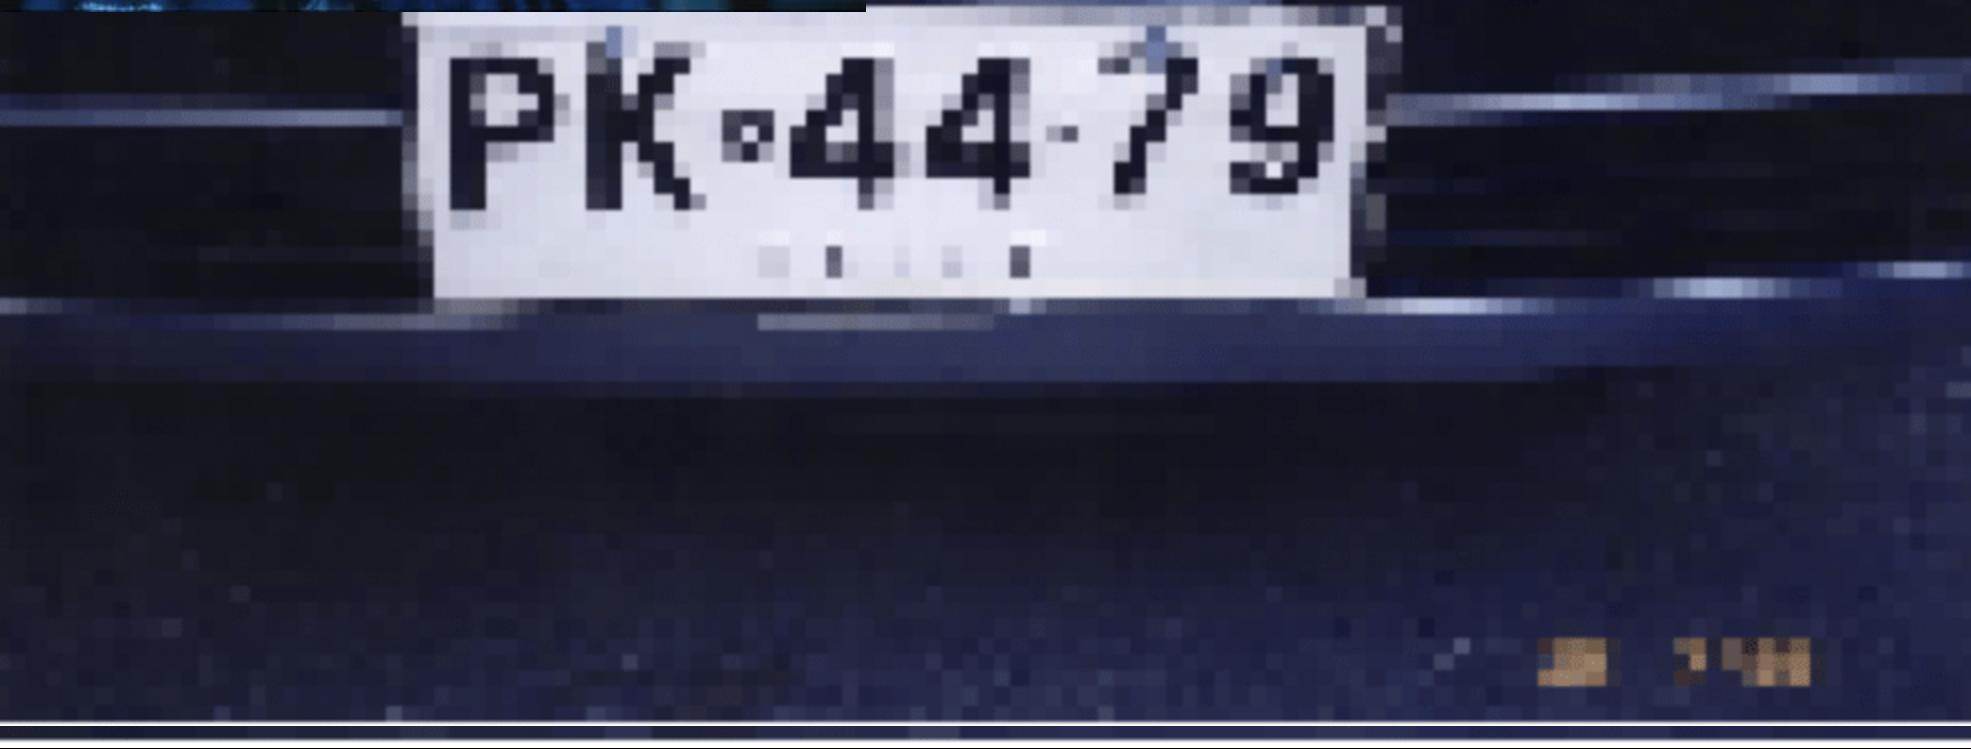
Super resolution in society and astronomy







CAN YOU
ENHANCE THAT



PK-44-79



CAN YOU
ENHANCE THAT

PK-44-79
CHILE

25 7 '99





remember the timae when simulations drove astronomy...

Experiment driven

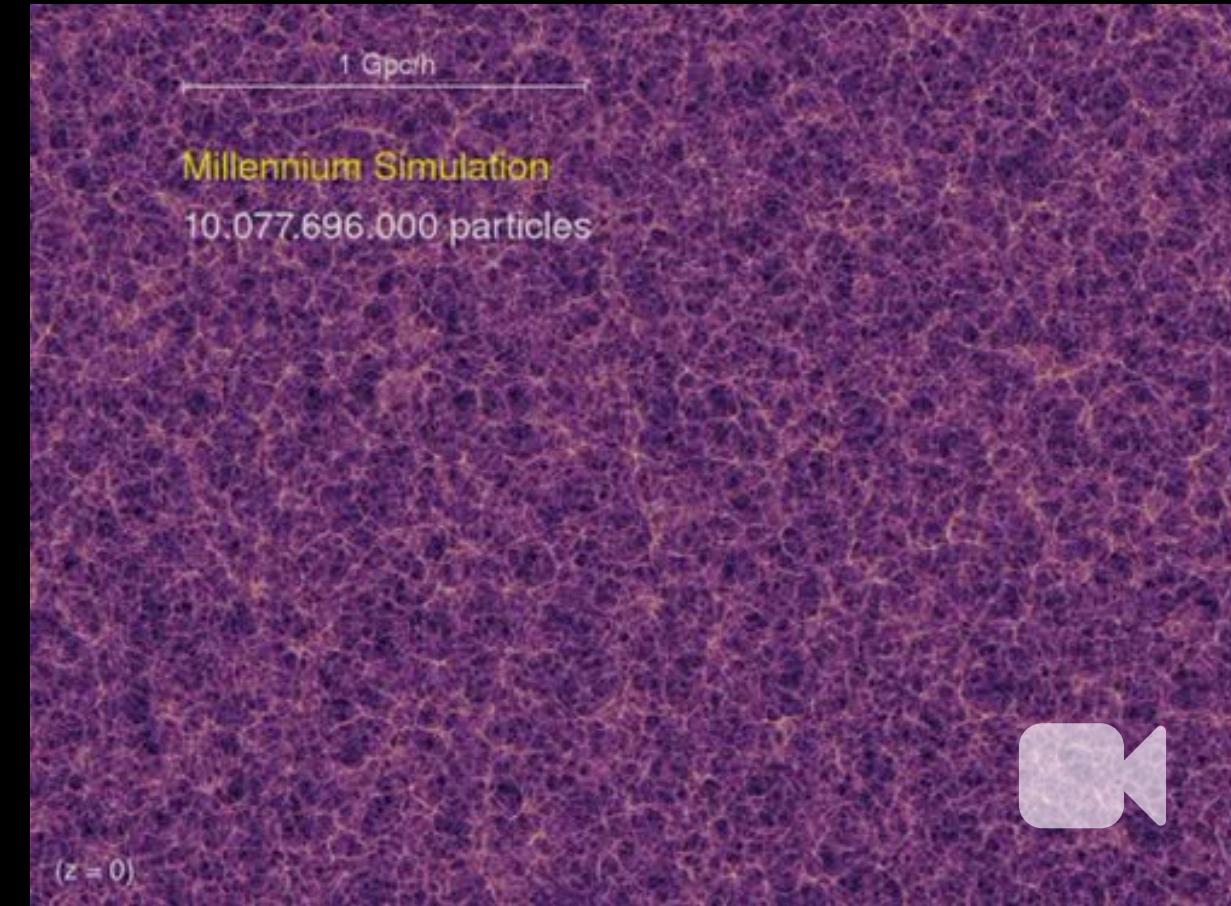
Theory driven | Falsifiability

Simulations | Probabilistic inference | Computation

The Millennium Run used more than 10^{10} particles to trace the evolution of the matter distribution in a cubic region of the Universe $500/h$ Mpc on a side (~over 2 billion light-years on a side), and has a spatial resolution of $5/h$ kpc. ~20M galaxies.

350 000 processor hours of CPU time, or 28 days of wall-clock time. Springel+2005

<https://wwwmpa.mpa-garching.mpg.de/galform/virgo/millennium>

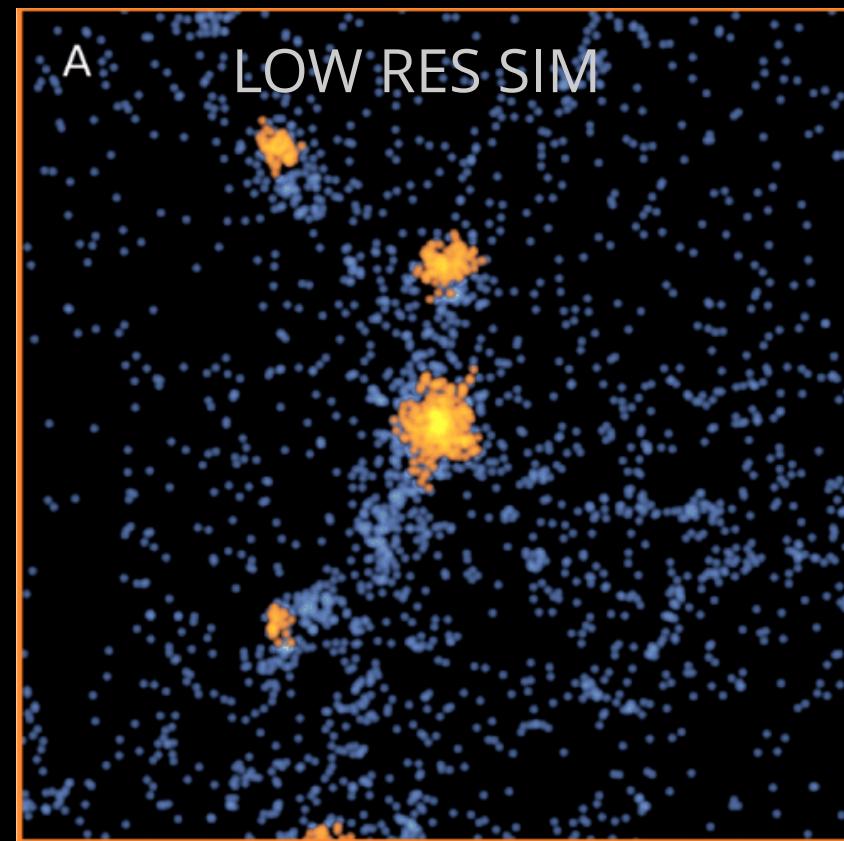


AI-assisted superresolution cosmological simulations

Yin Li+2021

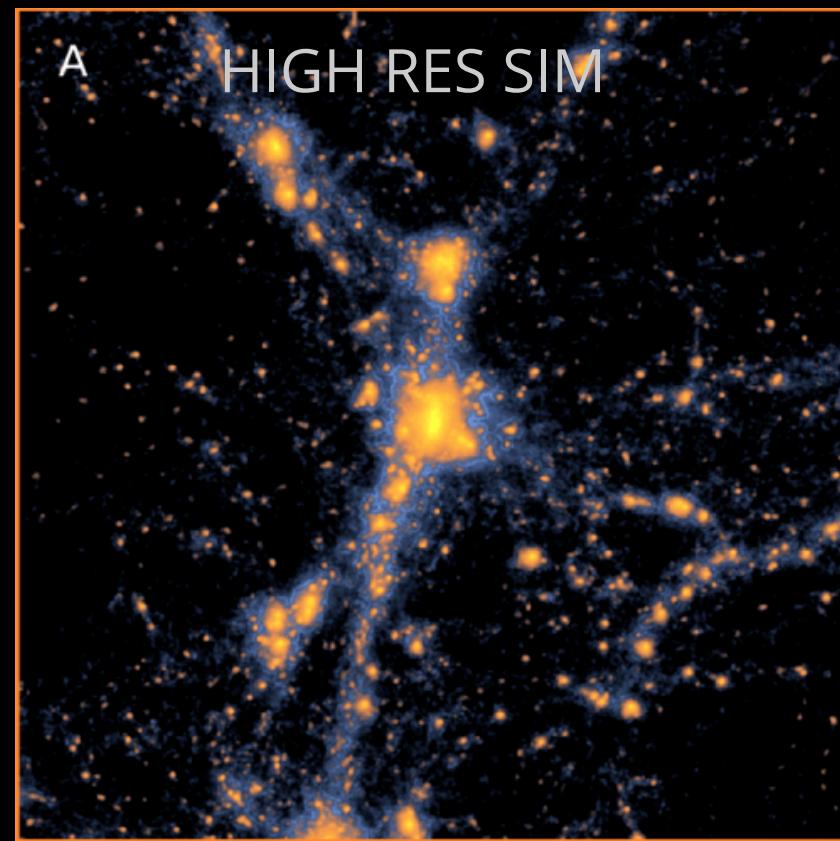
A

LOW RES SIM



A

HIGH RES SIM

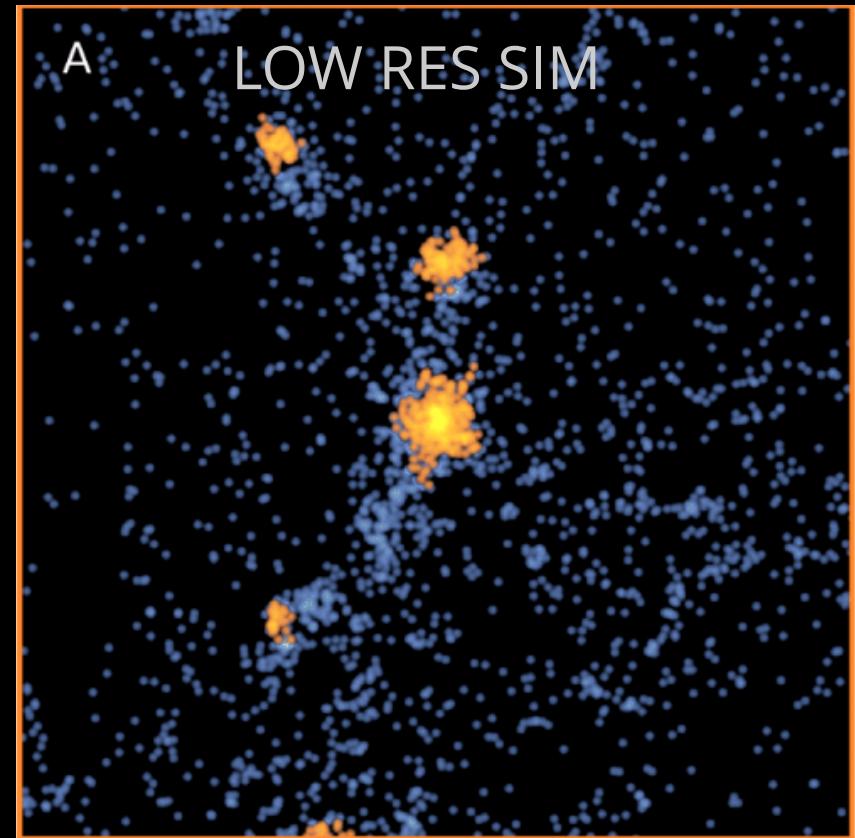


AI-assisted superresolution cosmological simulations

Yin Li+2021

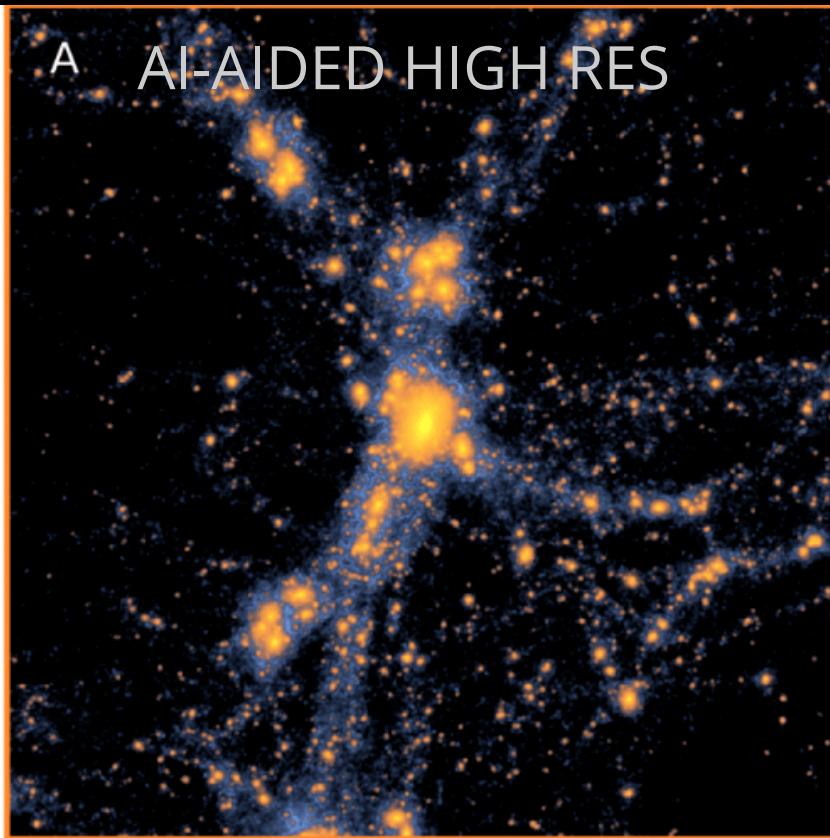
A

LOW RES SIM



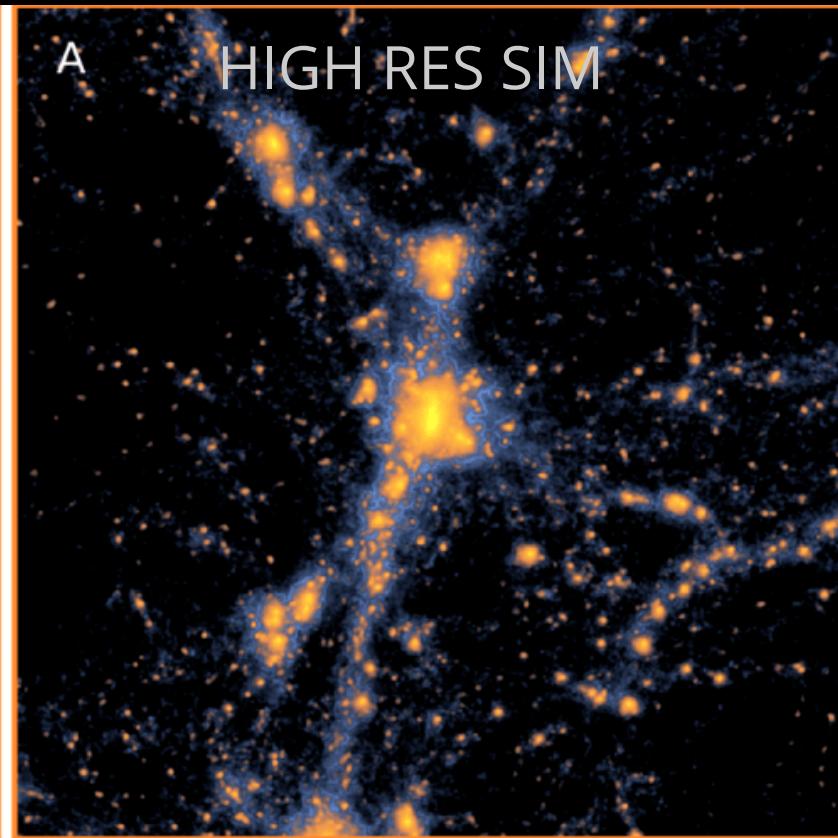
A

AI-AIDED HIGH RES



A

HIGH RES SIM

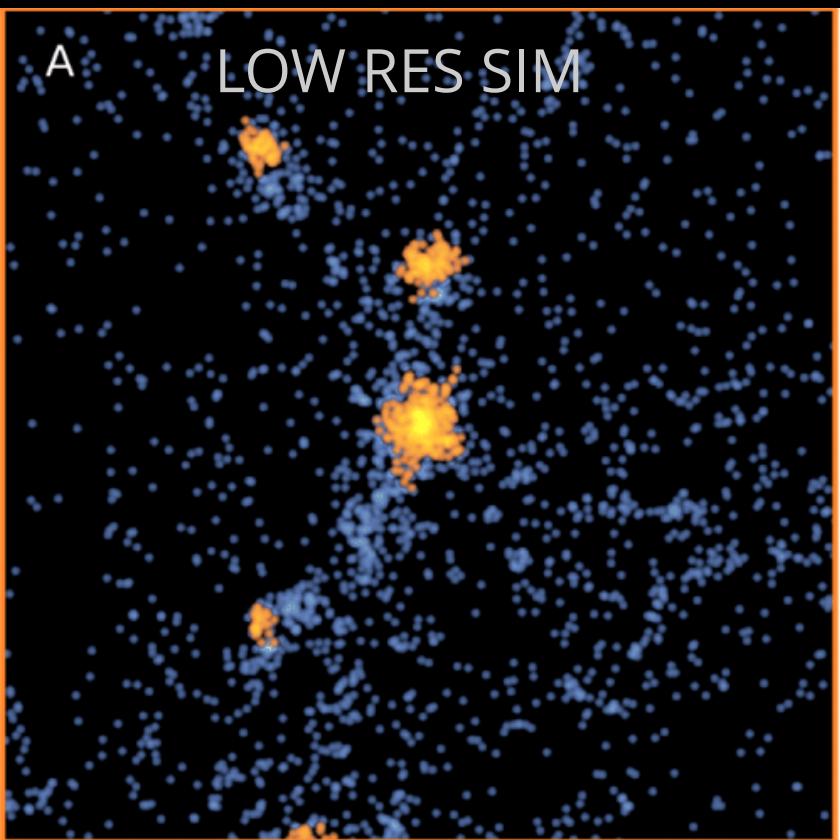


AI-assisted superresolution cosmological simulations

Yin Li+2021

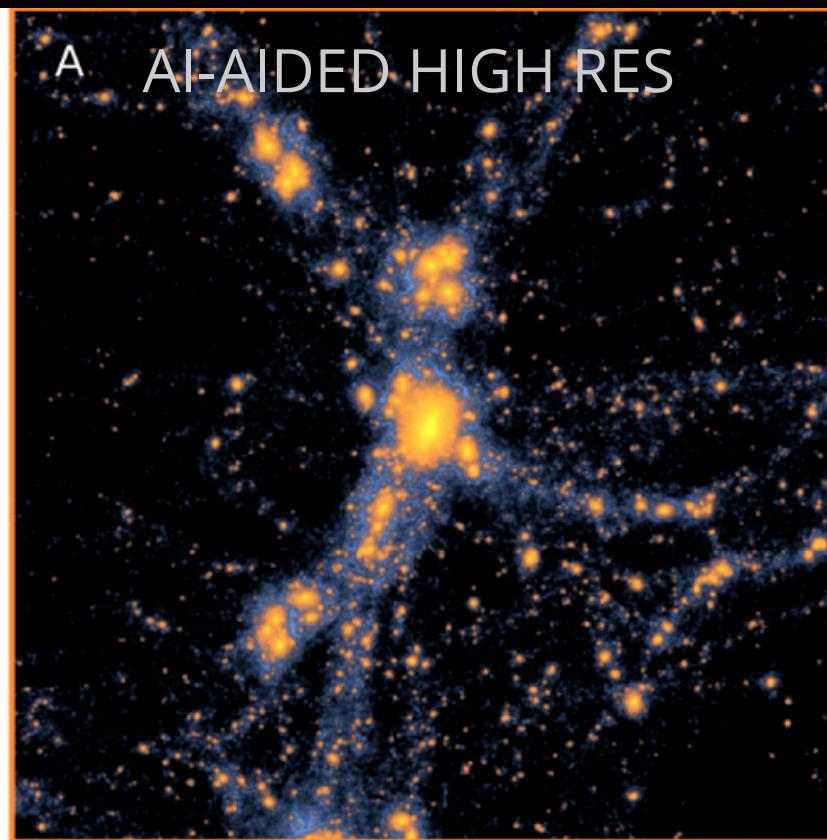
INPUT

A LOW RES SIM



OUTPUT

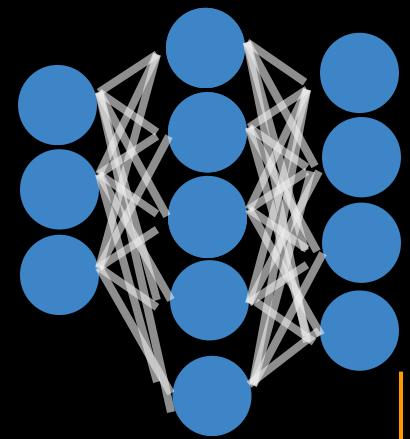
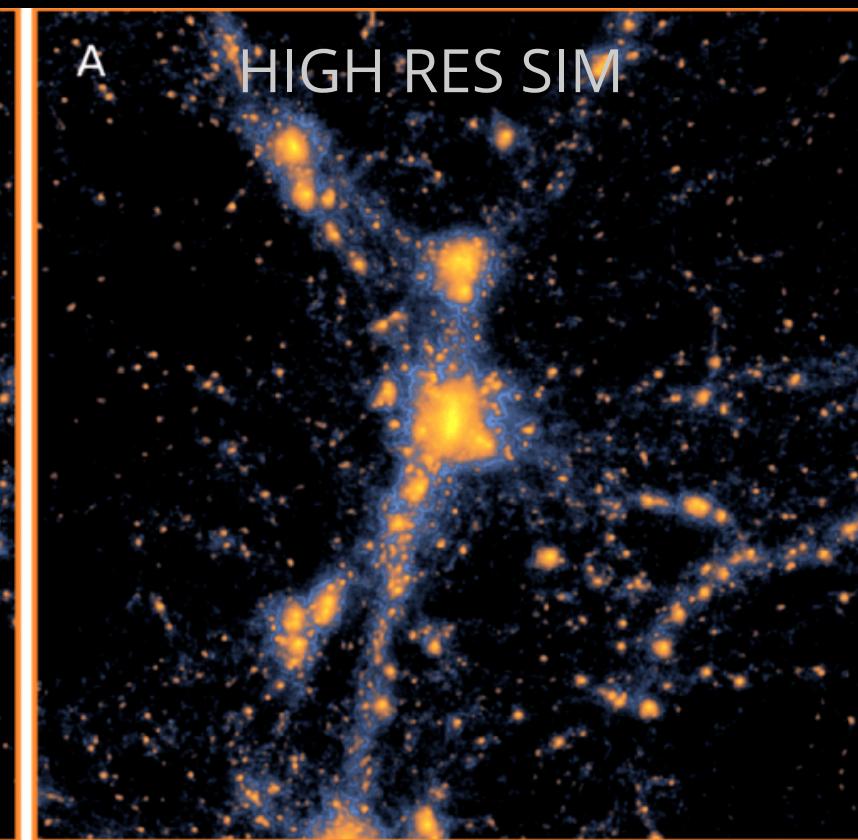
A AI-AIDED HIGH RES



$$\text{loss} = D(\text{OUTPUT-TARGET})$$

TARGET

A HIGH RES SIM

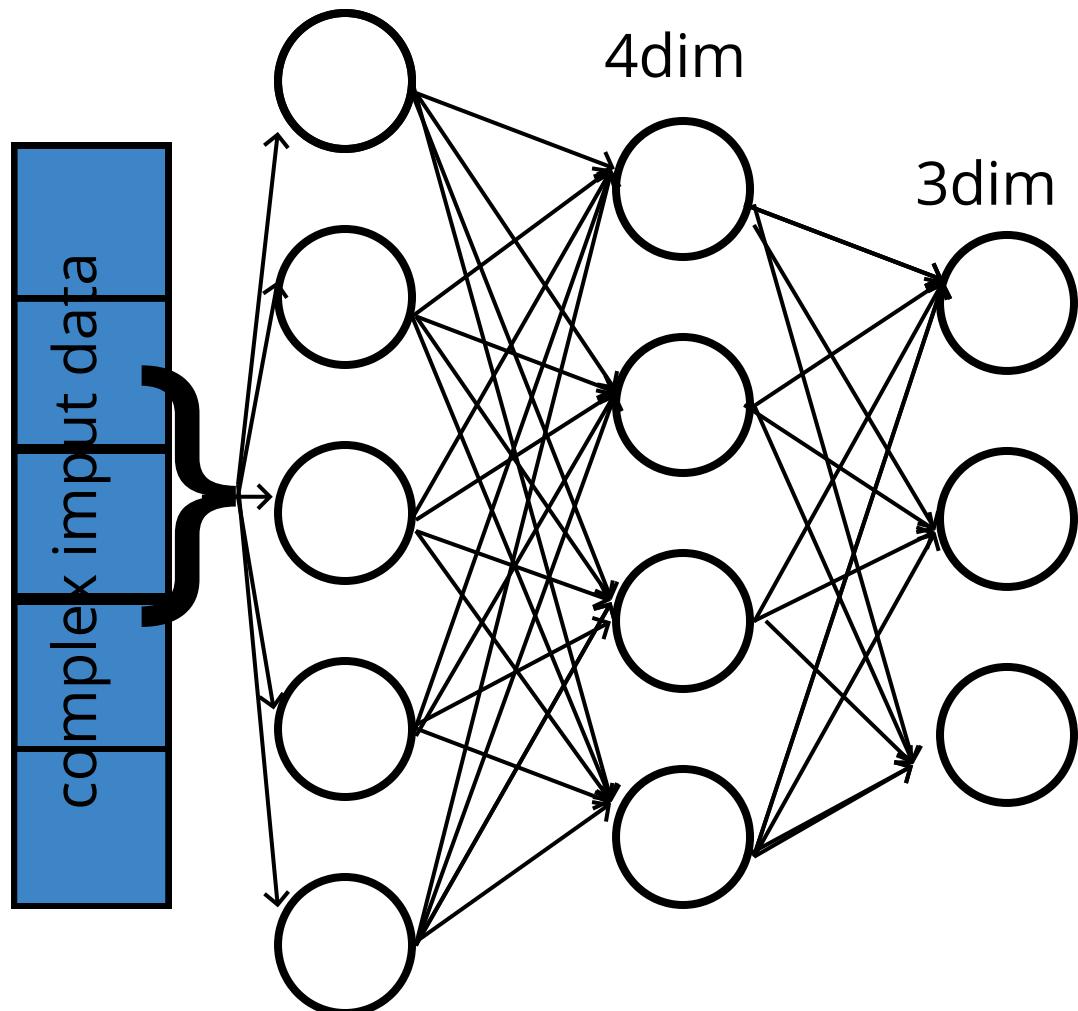


4/6 AUTOENCODERS

Unsupervised learning with Neural Networks

What do NN do? approximate complex functions with series of modified linear functions

5dim representation



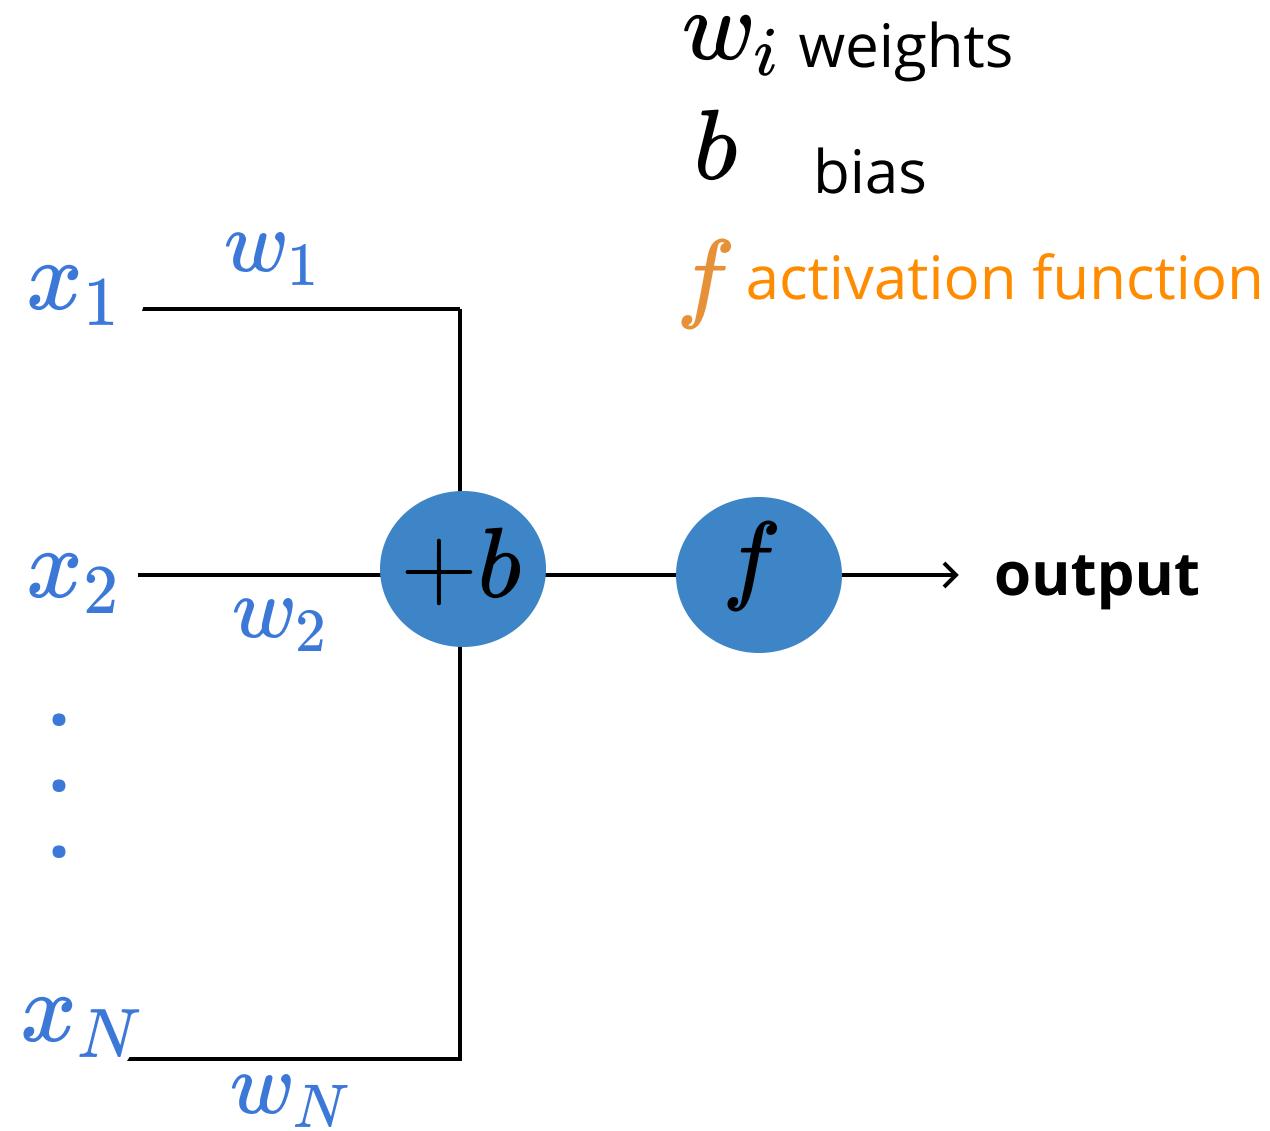
perceptrons

Perceptrons are **linear classifiers**:
makes its predictions based on a
linear predictor function
combining a set of weights
(=parameters) with the feature vector.

$$y = wx + b$$

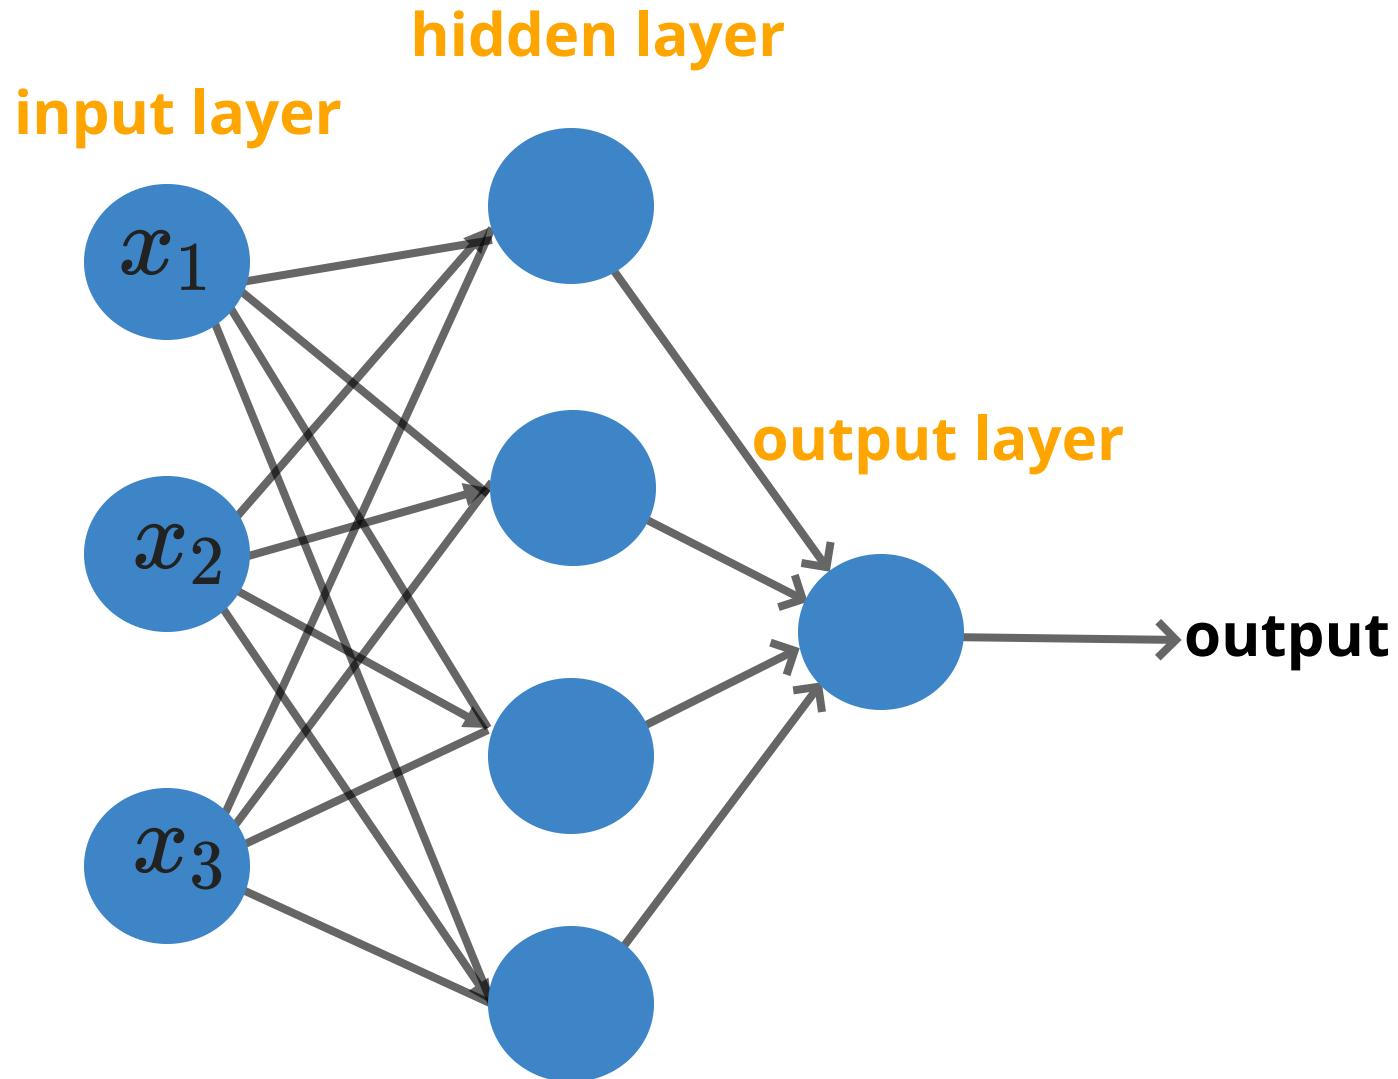
$$y = \sum_i w_i x_i + b$$

$$y = f(\sum_i w_i x_i + b)$$



multilayer perceptron

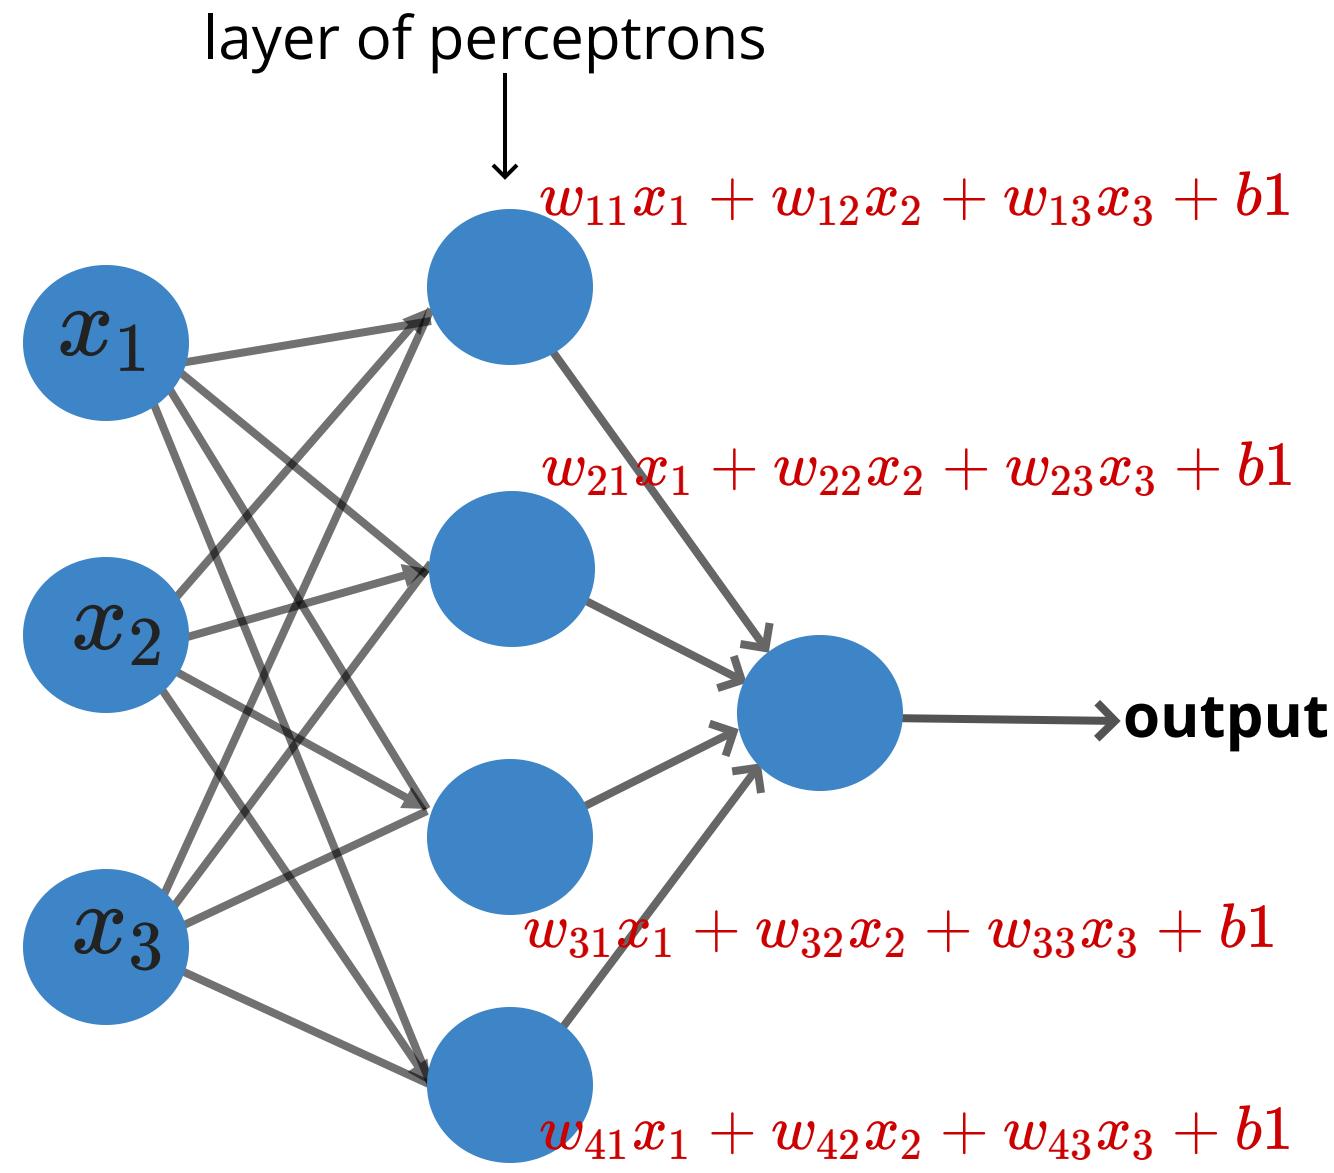
1970: multilayer
perceptron architecture



Fully connected: all nodes go to
all nodes of the next layer.

multilayer perceptron

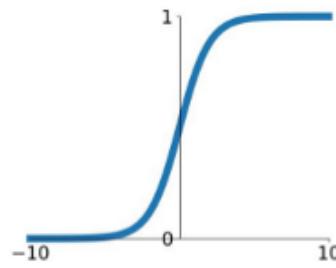
Fully connected: all nodes go to all nodes of the next layer.



activation functions

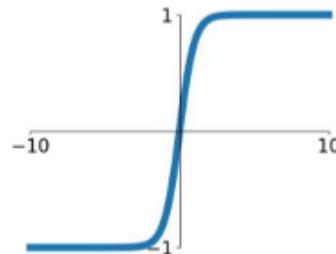
Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



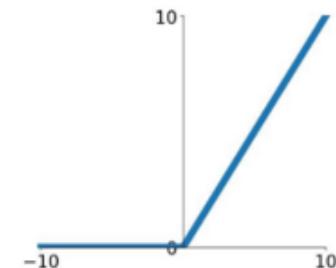
tanh

$$\tanh(x)$$



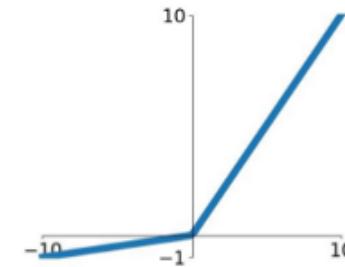
ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$

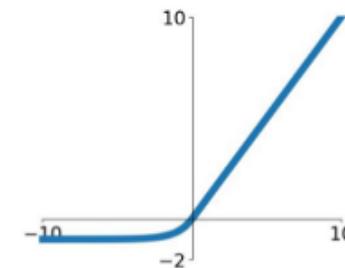


Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

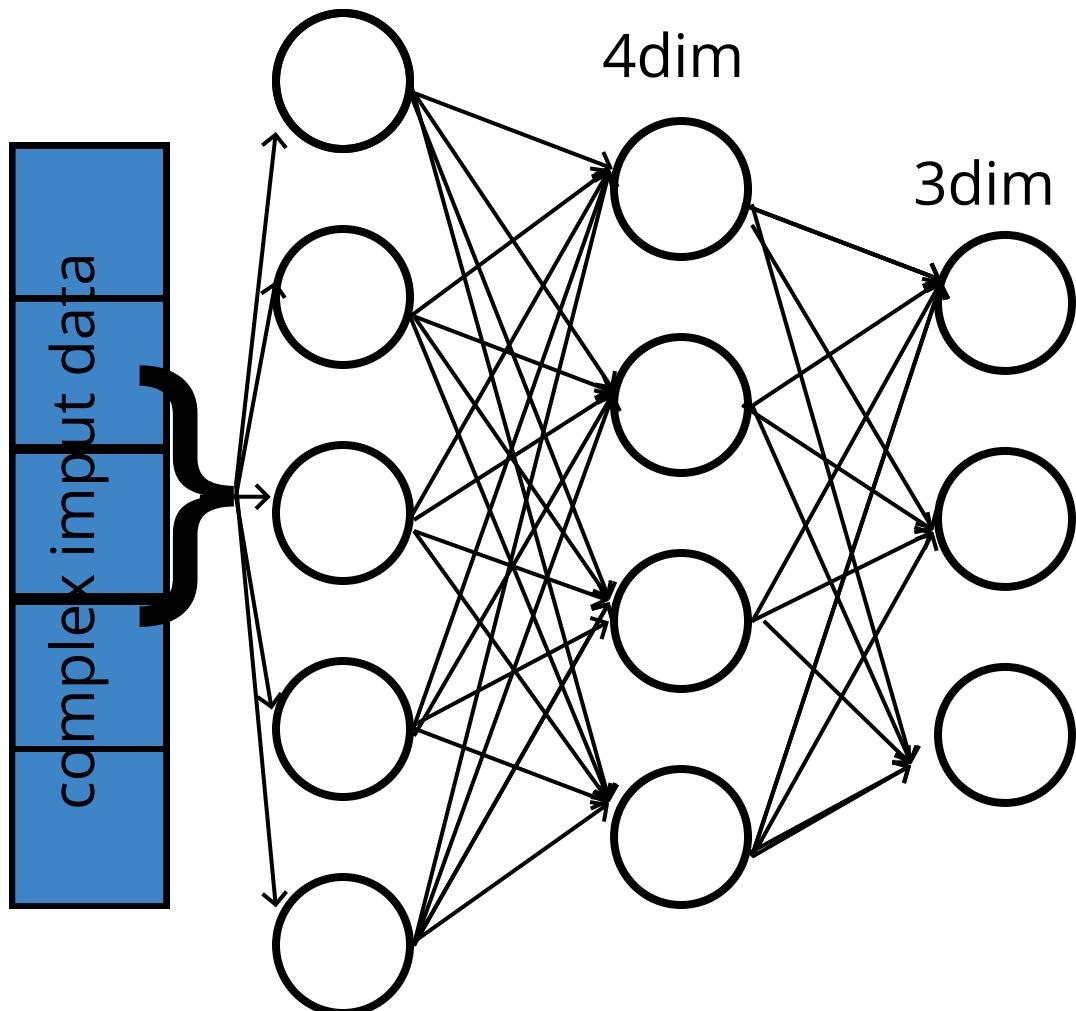


Unsupervised learning with Neural Networks

What do NN do? approximate complex functions with series of modified linear functions

To do that they extract information from the data: **each layer of the DNN produces a representation of the data a "latent representation"**

5dim representation



Unsupervised learning with Neural Networks

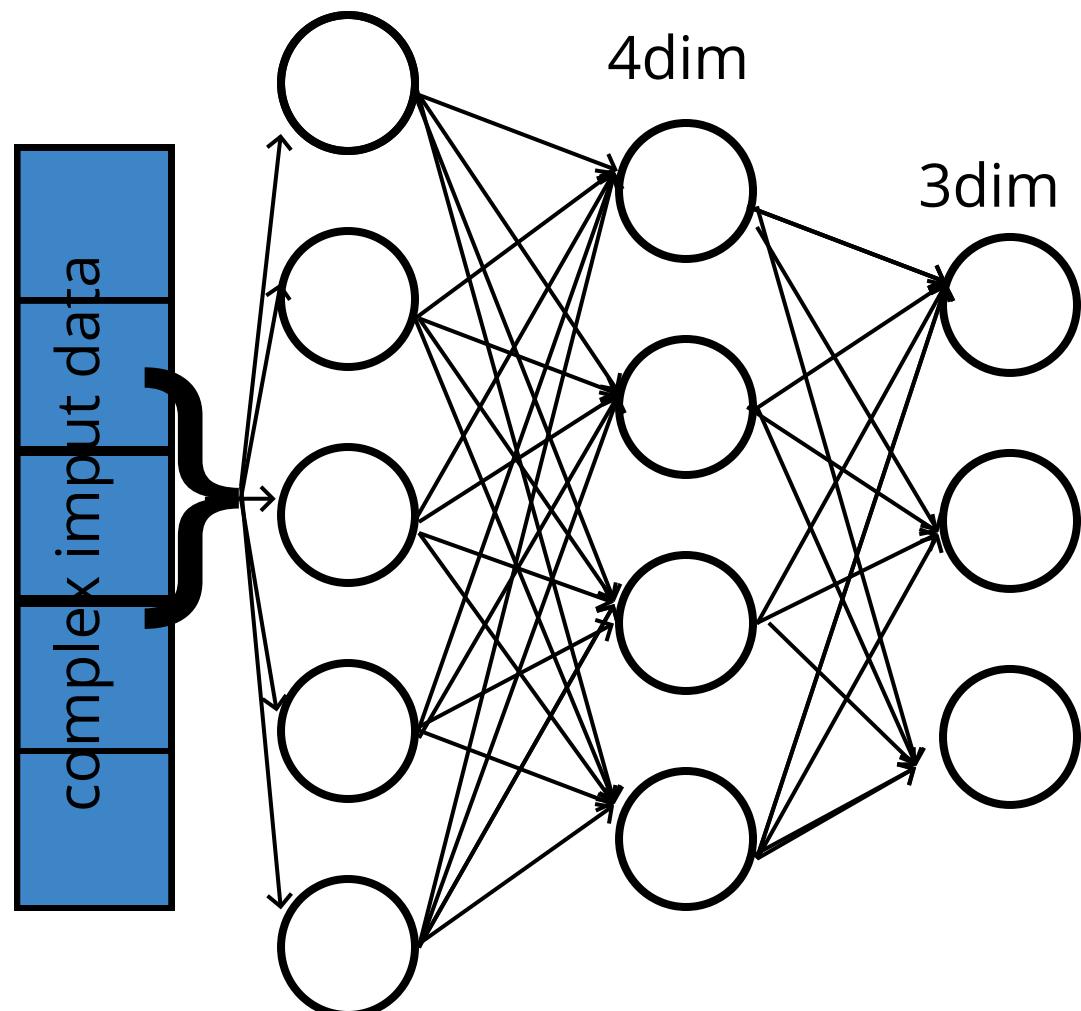
What do NN do? approximate complex functions with series of modified linear functions

To do that they extract information from the data: **each layer of the DNN produces a representation of the data a "latent representation"**

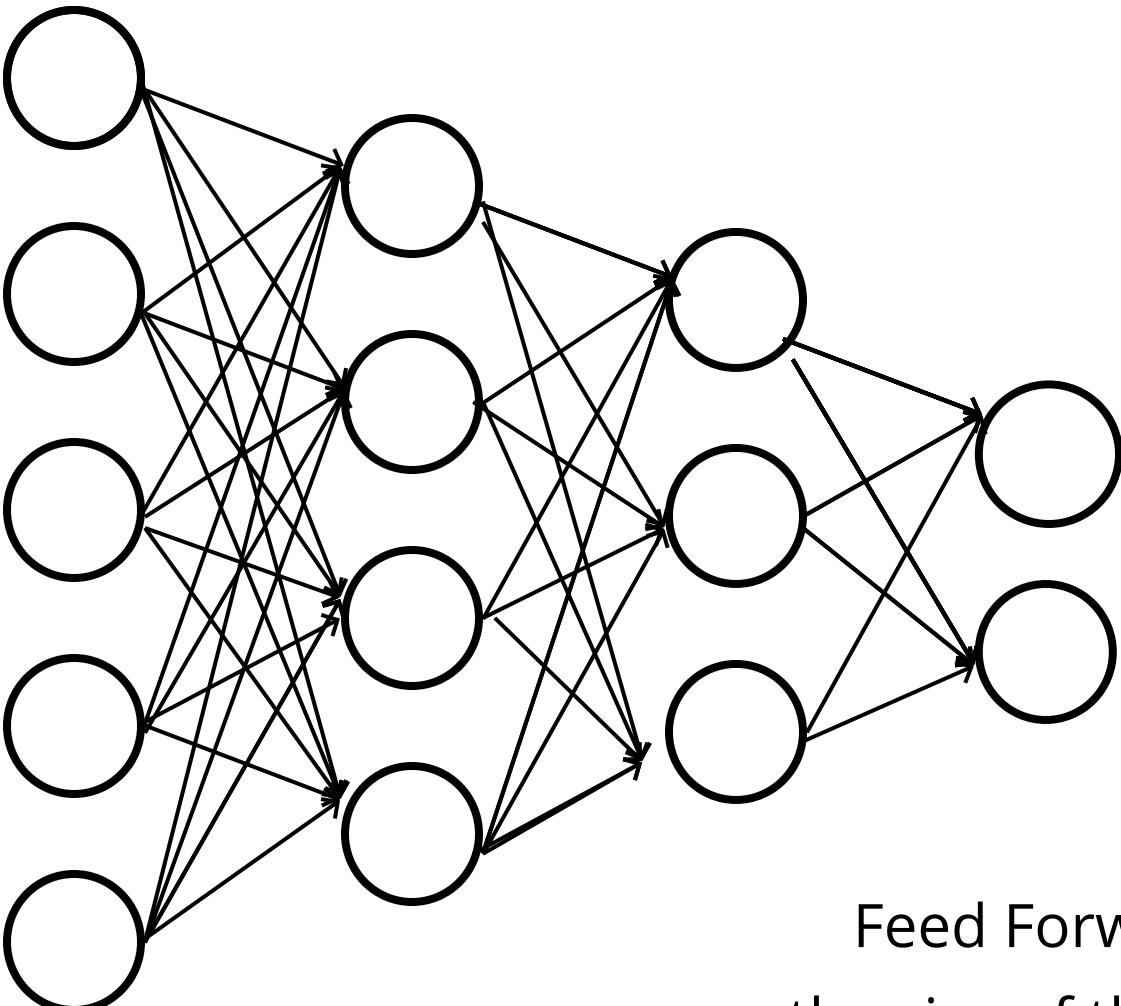
The dimensionality of that latent representation is determined by the size of the layer

.... so if my layers are smaller what I have is a compact representation of the data

5dim representation

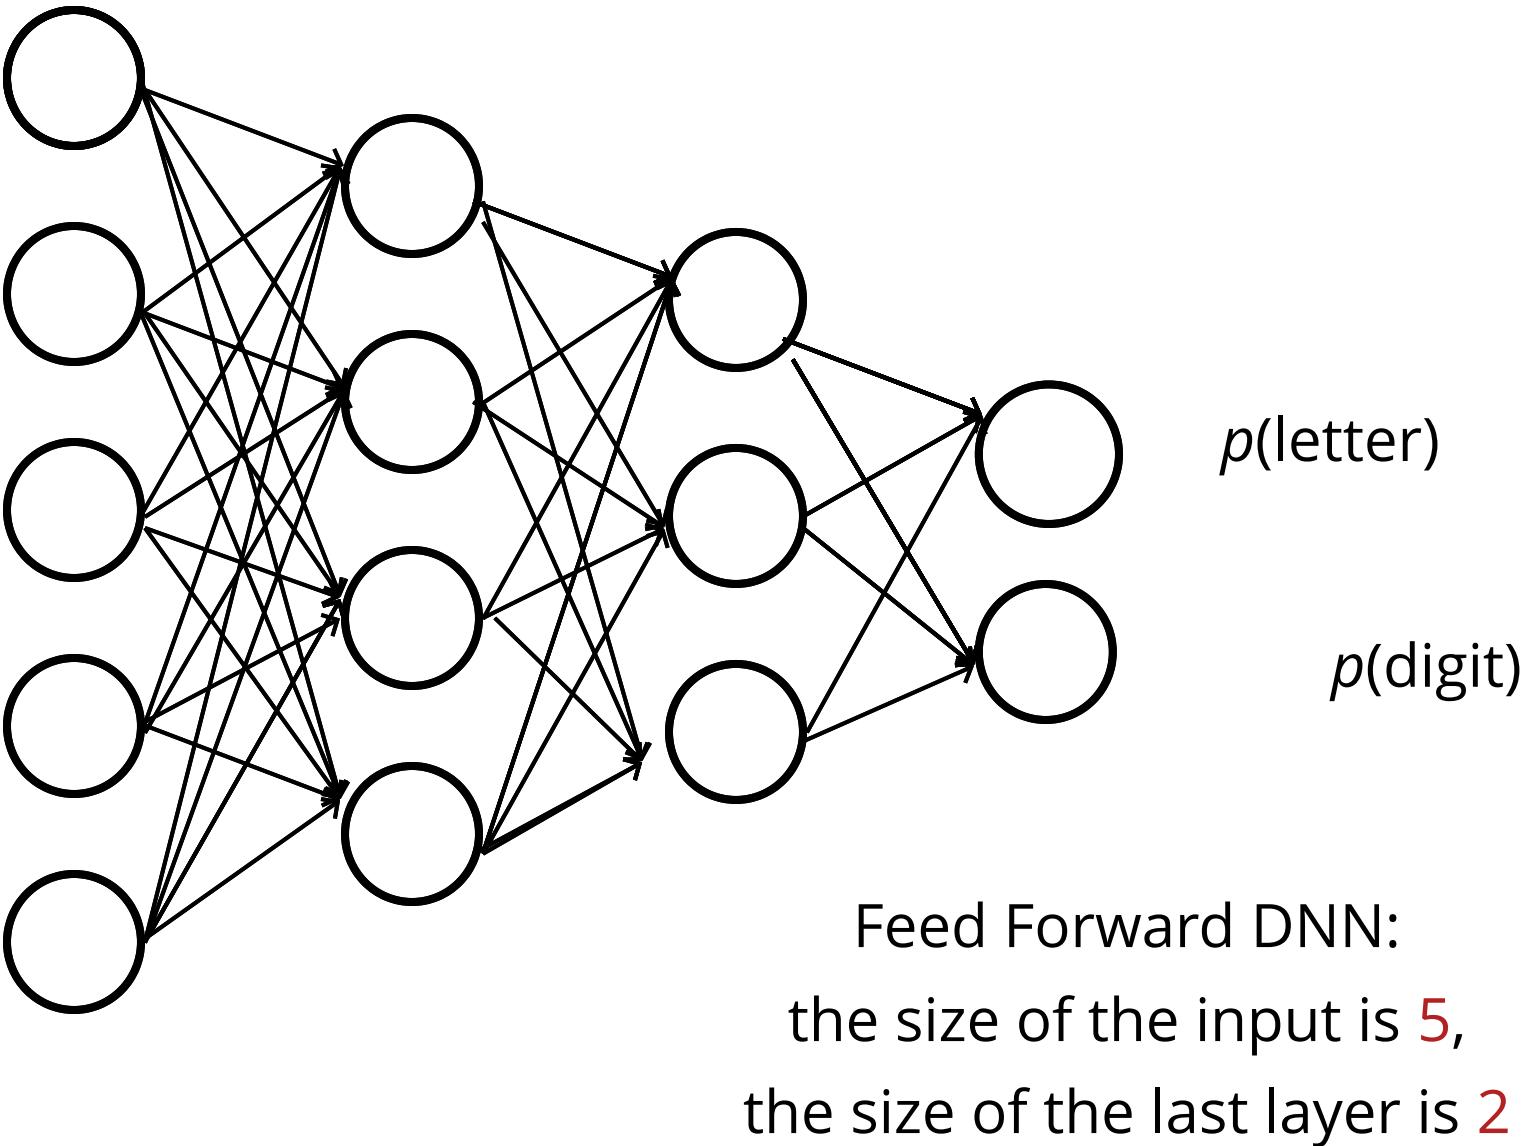


Autoencoder Architecture

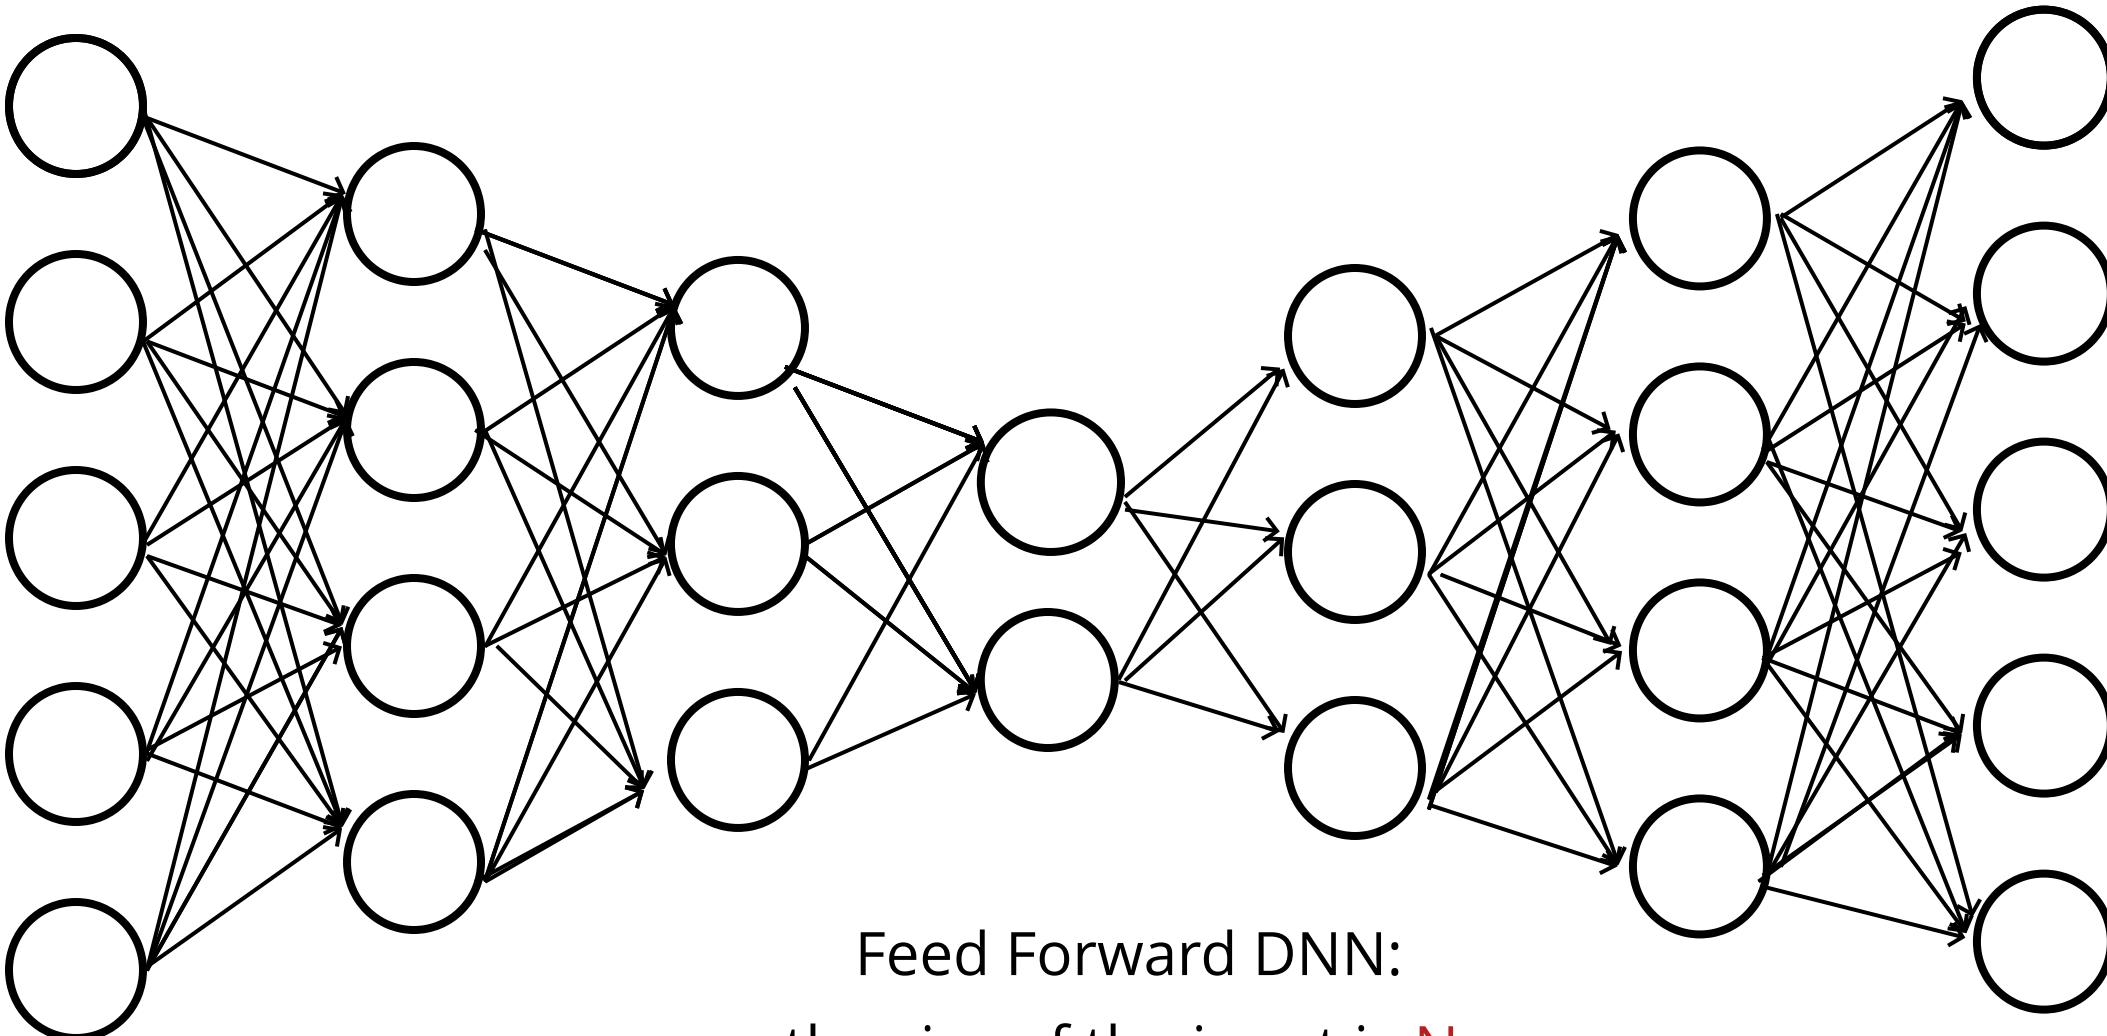


Feed Forward DNN:
the size of the input is 5,
the size of the last layer is 2

Autoencoder Architecture



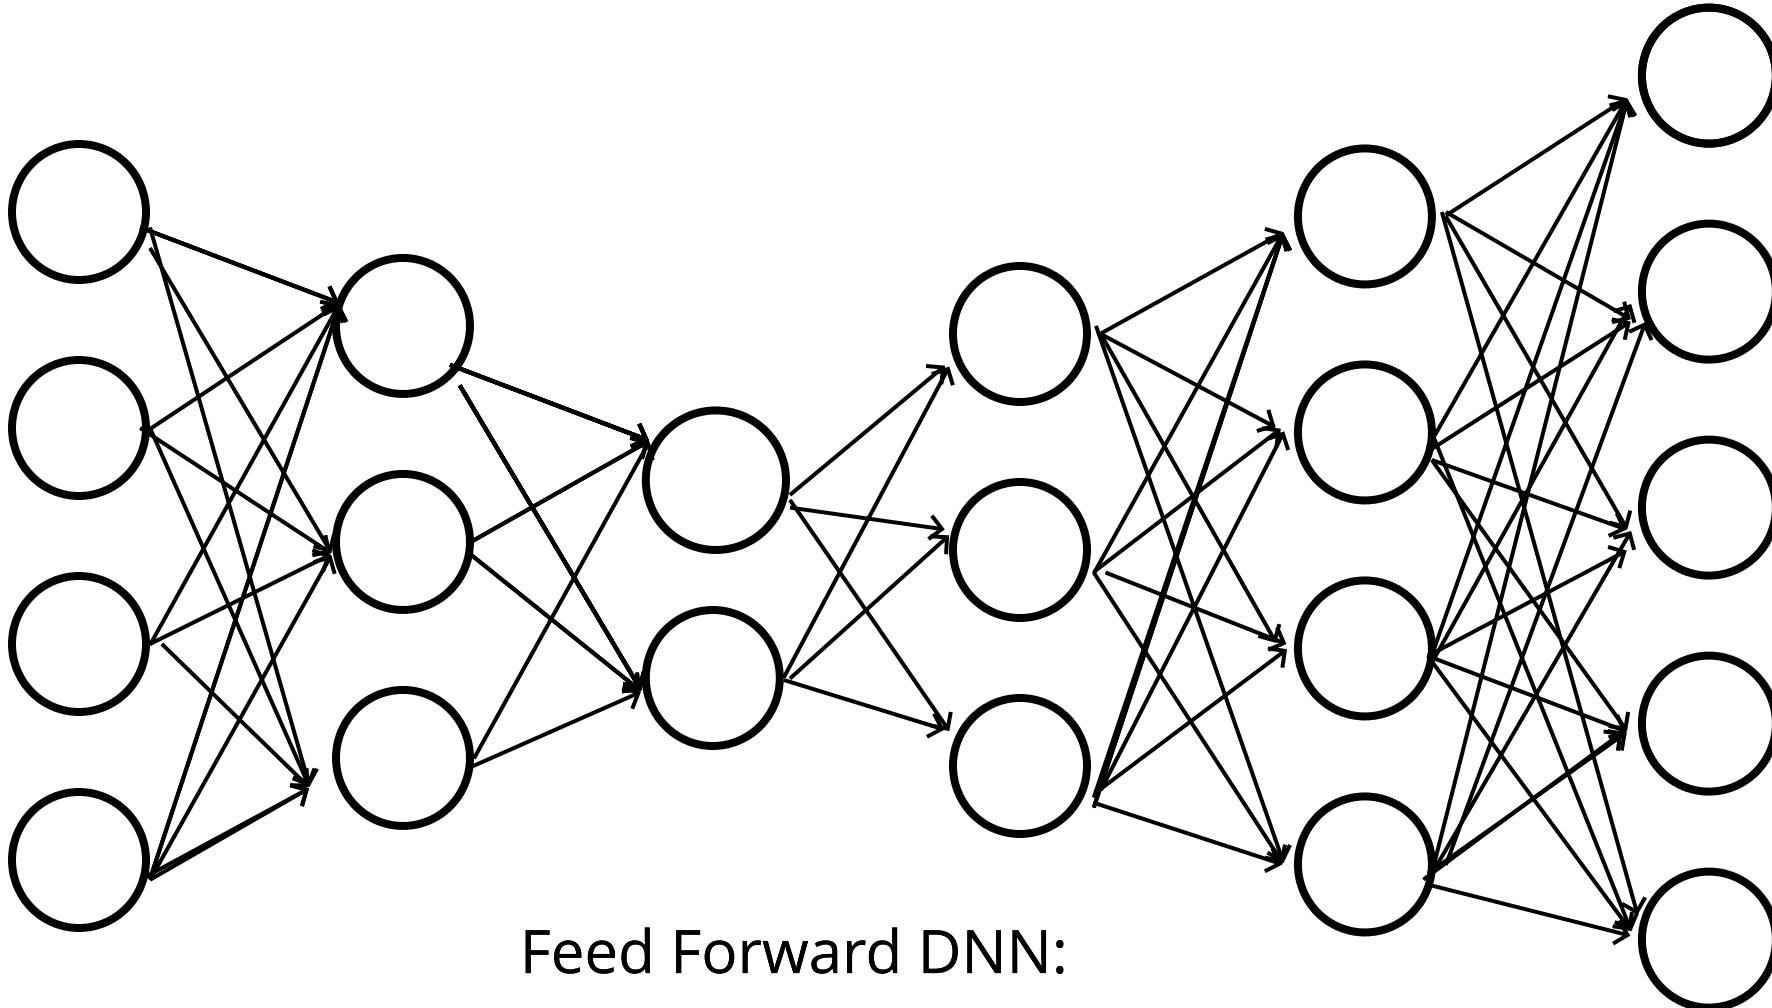
Autoencoder Architecture



Feed Forward DNN:
the size of the input is **N**,
the size of the last layer is **N**



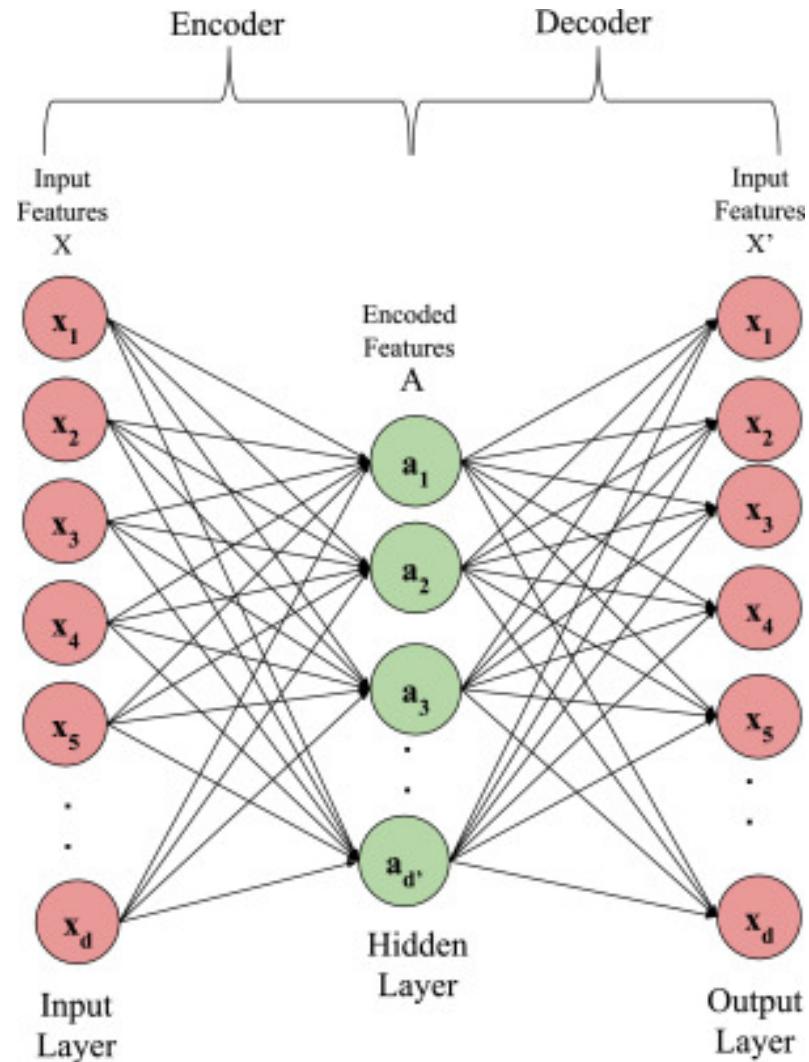
Autoencoder Architecture



Feed Forward DNN:
the size of the input is N ,
the size of the last layer is N

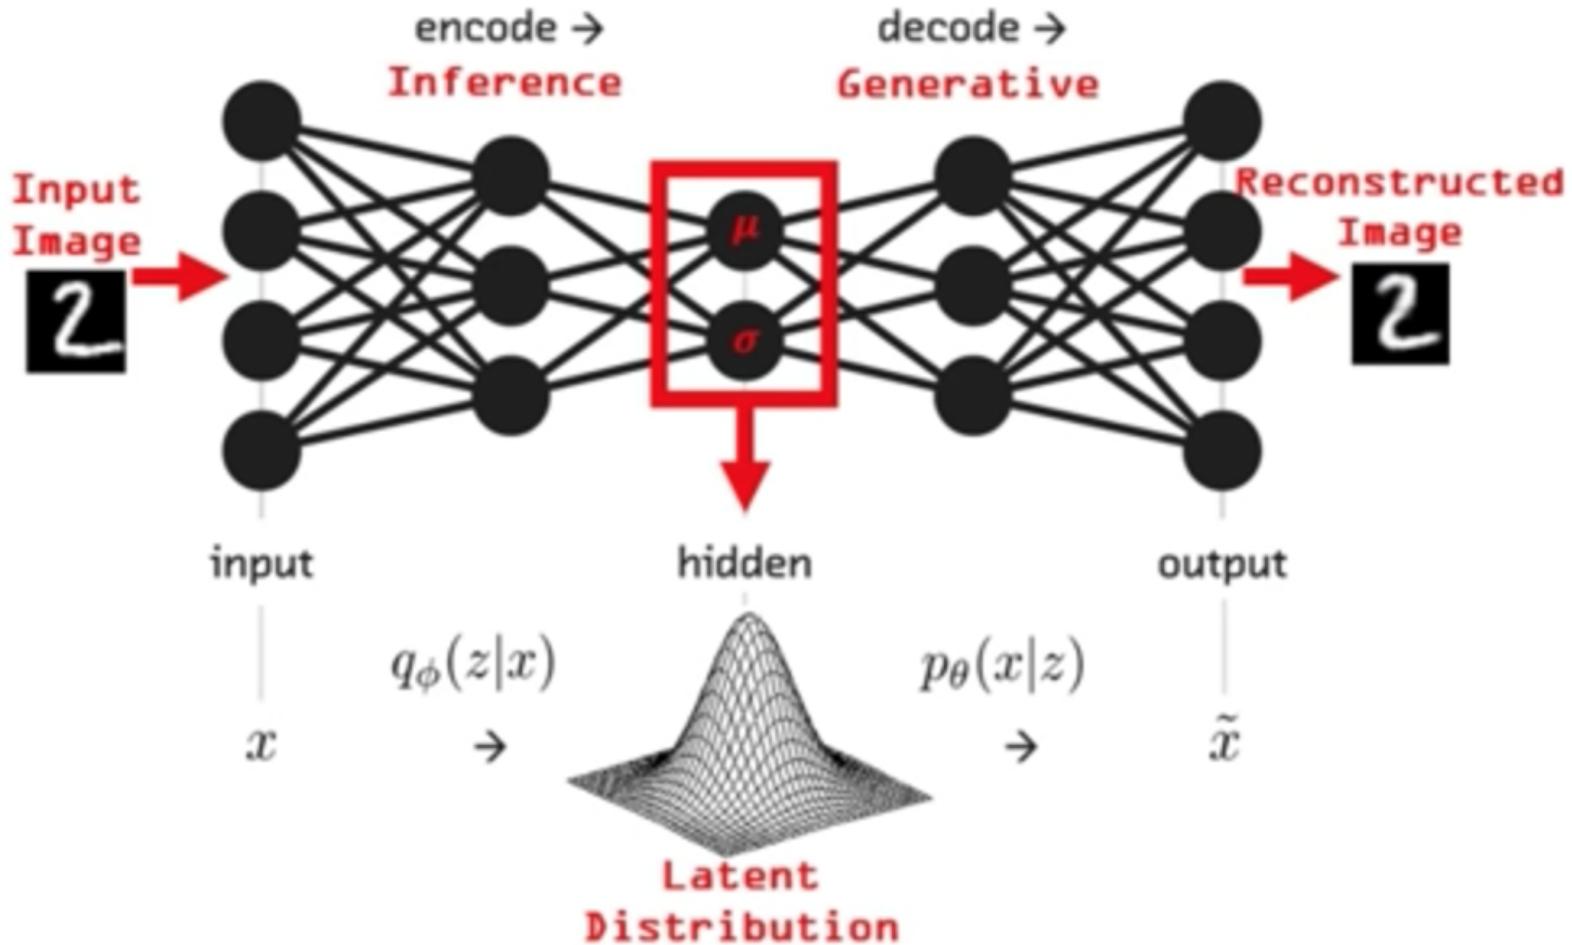
Autoencoder Architecture

https://link.springer.com/chapter/10.1007/978-981-13-6661-1_3



- **Encoder:** outputs a lower dimensional representation \mathbf{z} of the data \mathbf{x} (similar to PCA, tSNE...)
- **Decoder:** Learns how to reconstruct \mathbf{x} given \mathbf{z} : learns $p(\mathbf{x} | \mathbf{z})$

Variational Autoencoder Architecture



The hidden representation is assumed (forced) to be Gaussian. The bottle-neck layer represents the mean and standard deviation of that distribution

Building a DNN with keras and tensorflow

Trivial to build, but the devil is in the details!

Building a DNN with keras and tensorflow

Trivial to build, but the devil is in the details!

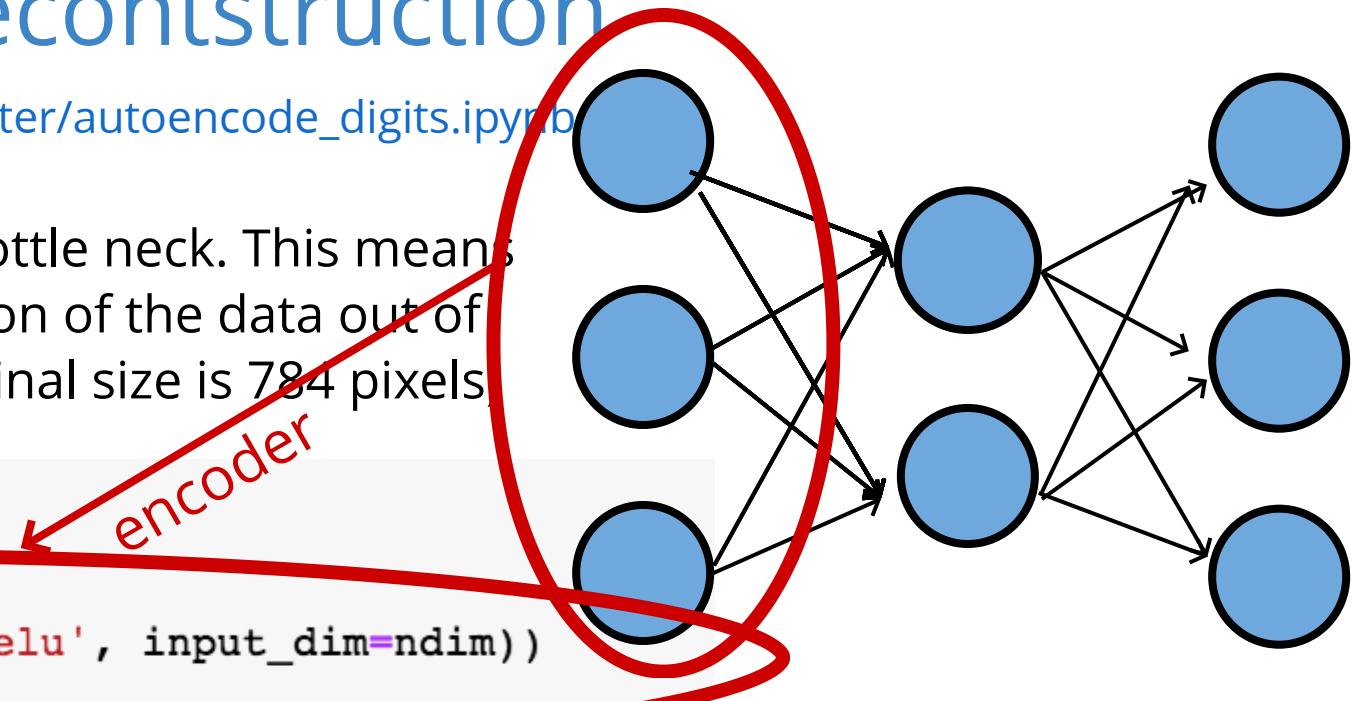
```
1 from keras.models import Sequential
2 #can upload pretrained models from keras.models
3 from keras.layers import Dense, Conv2D, MaxPooling2D
4 #create model
5 model = Sequential()
6
7
8 #create the model architecture by adding model layers
9 model.add(Dense(10, activation='relu', input_shape=(n_cols,)))
10 model.add(Dense(10, activation='relu'))
11 model.add(Dense(1))
12
13 #need to choose the loss function, metric, optimization scheme
14 model.compile(optimizer='adam', loss='mean_squared_error')
15
16 #need to learn what to look for - always plot the loss function!
17 model.fit(x_train, y_train, validation_data=(x_test, y_test),
18            epochs=20, batch_size=100, verbose=1)
19 #note that the model allows to give a validation test,
20 #this is for a 3fold cross valiation: train-validate-test
21 #model.evaluate
```

Building a DNN with keras and tensorflow autoencoder for image reconstruction

https://github.com/fedhere/MLTSA_FBianco/blob/master/autoencode_digits.ipynb

This autoencoder model has a 64-neuron bottle neck. This means it will generate a compressed representation of the data out of that layer which is 16-dimensional (the original size is 784 pixels).

```
model_digits64 = Sequential()
## encoder
# input layer and the output size
model_digits64.add(Dense(128, activation='relu', input_dim=ndim))
#compression layer
model_digits64.add(Dense(64, activation='relu'))
## deencoder
#decompression layer, same size as in the encoder
model_digits64.add(Dense(128, activation='relu'))
#output layer, same size as input
model_digits64.add(Dense(ndim, activation='linear'))
```

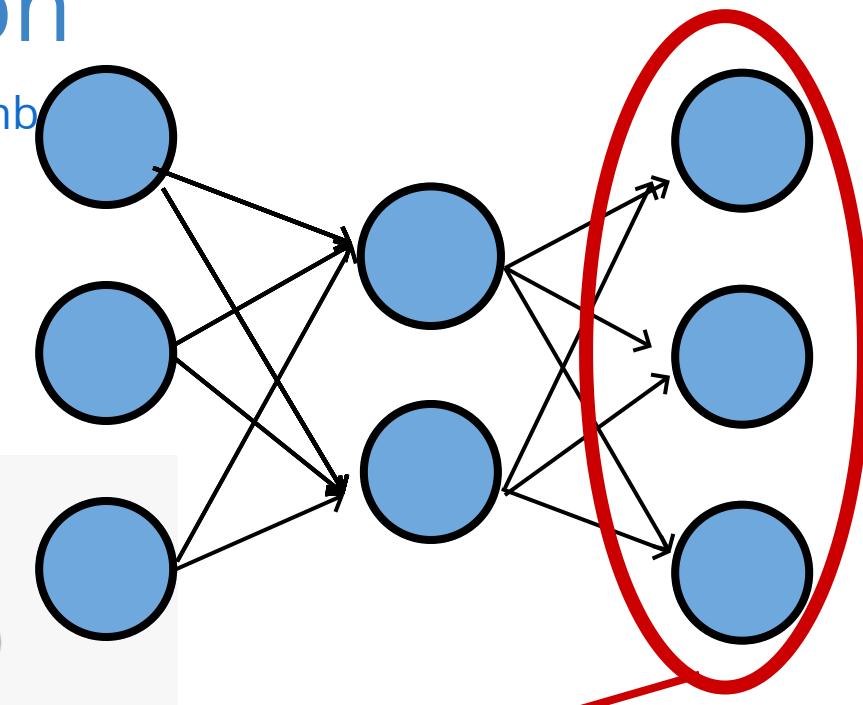


Building a DNN with keras and tensorflow autoencoder for image reconstruction

https://github.com/fedhere/MLTSA_FBianco/blob/master/autoencode_digits.ipynb

This autoencoder model has a 64-neuron bottle neck. This means it will generate a compressed representation of the data out of that layer which is 16-dimensional (the original size is 784 pixels)

```
model_digits64 = Sequential()
## encoder
# input layer and the output size
model_digits64.add(Dense(128, activation='relu', input_dim=ndim))
#compression layer
model_digits64.add(Dense(64, activation='relu'))
## deencoder
#decompression layer, same size as in the encoder
model_digits64.add(Dense(128, activation='relu'))
#output layer, same size as input
model_digits64.add(Dense(ndim, activation='linear'))
```



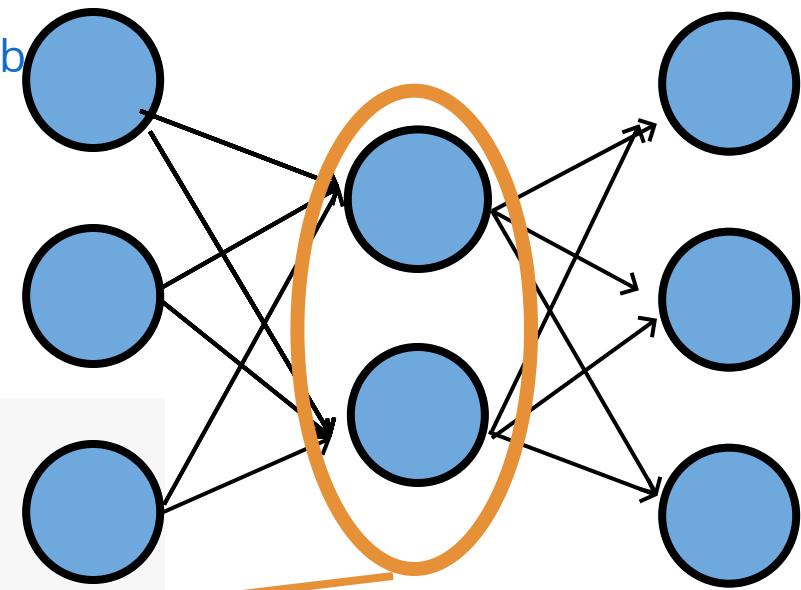
decoder

Building a DNN with keras and tensorflow autoencoder for image reconstruction

https://github.com/fedhere/MLTSA_FBianco/blob/master/autoencode_digits.ipynb

This autoencoder model has a 64-neuron bottle neck. This means it will generate a compressed representation of the data out of that layer which is 16-dimensional (the original size is 784 pixels)

```
model_digits64 = Sequential()
## encoder
# input layer and the output size
model_digits64.add(Dense(128, activation='relu', input_dim=ndim))
#compression layer
model_digits64.add(Dense(64, activation='relu')) ← bottle neck
## deencoder
#decompression layer, same size as in the encoder
model_digits64.add(Dense(128, activation='relu'))
#output layer, same size as input
model_digits64.add(Dense(ndim, activation='linear'))
```



Building a DNN with keras and tensorflow autoencoder for image reconstruction

https://github.com/fedhere/MLTSA_FBianco/blob/master/autoencode_digits.ipynb

```
# choose the optimizer and loss appropriately!
model_digits64.compile(optimizer="adadelta", loss="mean_squared_error")
```

```
print(model_digits64.summary())
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 128)	100480
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 128)	8320
dense_4 (Dense)	(None, 784)	101136

Total params: 218,192

Trainable params: 218,192

Non-trainable params: 0

None

This simple model has 200K parameters!

My original choice is to train it with "adadelta" with a mean squared loss function, all activation functions are relu, appropriate for a linear regression

Building a DNN with keras and tensorflow autoencoder for image reconstruction

https://github.com/fedhere/MLTSA_FBianco/blob/master/autoencode_digits.ipynb

3.0.1 regression

- loss='mean_squared_error' L2: default loss to use for regression problems. => linear activation function in output layer, one node out

alternatives: loss='mean_squared_logarithmic_error', 'mean_absolute_error' (which is L1 instead of L2)

3.0.2 binary classification

- loss='binary_crossentropy' => sigmoid activation function in output layer, one node out

alternatives: 'hinge'

3.0.3 multiclass classification

categorical encoded as numerical

- loss='categorical_crossentropy' => softmax n nodes out

onehot encoded categorical

- 'sparse_categorical_crossentropy' => softmax n nodes out
- 'kullback Leibler Divergence Loss' => probabilistic categorical classification; $\log(P/Q)$

What should I choose for the loss function and how does that relate to the activation function and optimization?

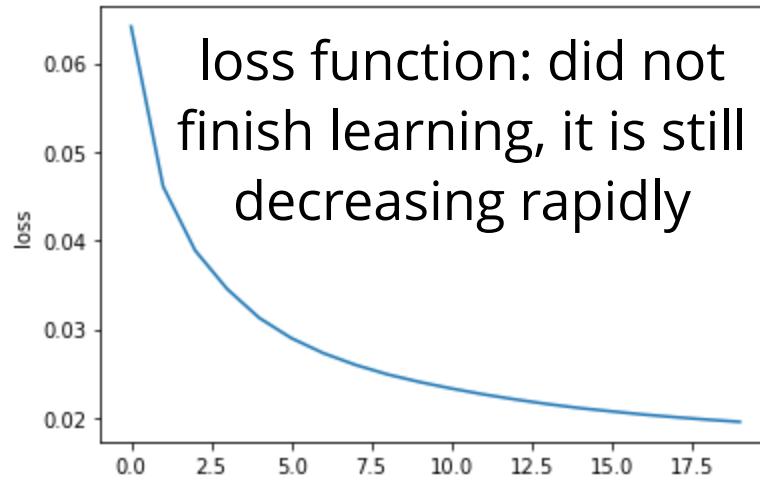
autoencoder for image reconstruction

https://github.com/fedhere/MLTSA_FBianco/blob/master/autoencode_digits.ipynb

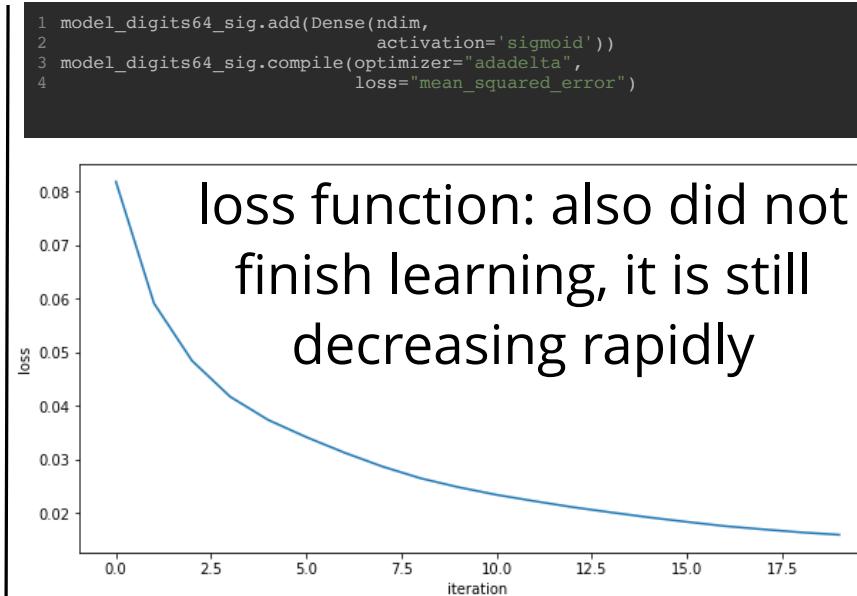
```
1 model_digits64.add(Dense(ndim,  
2                         activation='linear'))  
3 model_digits64_sig.compile(optimizer="adadelta",  
4                           loss="mean_squared_error")
```

```
1 model_digits64_sig.add(Dense(ndim,  
2                             activation='sigmoid'))  
3 model_digits64_sig.compile(optimizer="adadelta",  
4                           loss="mean_squared_error")
```

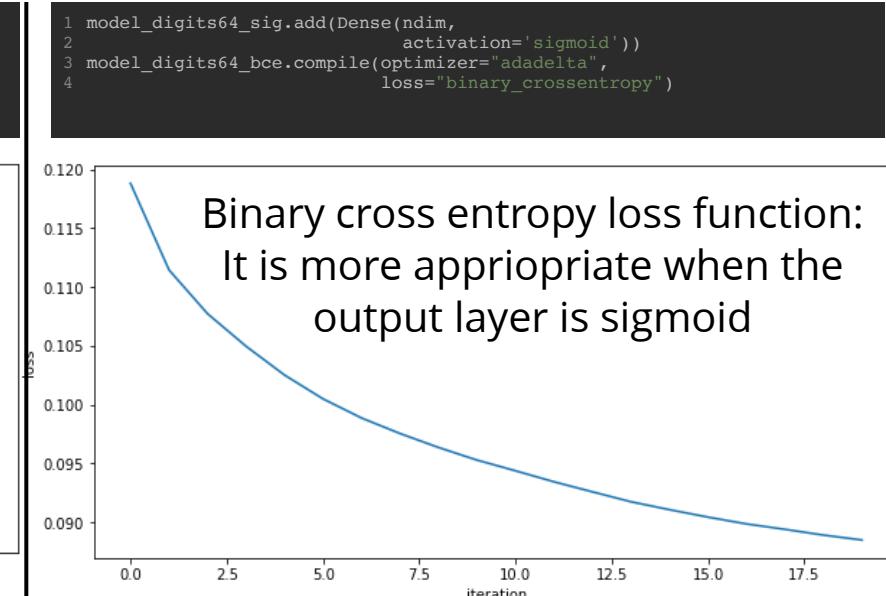
```
1 model_digits64_sig.add(Dense(ndim,  
2                             activation='sigmoid'))  
3 model_digits64_bce.compile(optimizer="adadelta",  
4                           loss="binary_crossentropy")
```



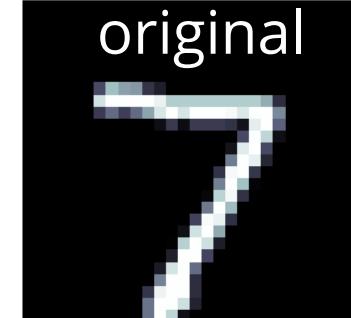
The predictions are far too detailed. While the input is not binary, it does not have a lot of details. Maybe approaching it as a binary problem (with a sigmoid and a binary cross entropy loss) will give better results



A sigmoid gives activation gives a much better result!



Even better results!



autoencoder for image reconstruction

https://github.com/fedhere/MLTSA_FBianco/blob/master/autoencode_digits.ipynb

A more ambitious model has a 16 neurons bottle neck:
we are trying to extract 16 numbers to reconstruct the
entire image! its pretty remarkable! those 16 number
are ***extracted features*** from the data

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_9 (Dense)	(None, 128)	100480
dense_10 (Dense)	(None, 64)	8256
dense_11 (Dense)	(None, 32)	2080
dense_12 (Dense)	(None, 16)	528
dense_13 (Dense)	(None, 32)	544
dense_14 (Dense)	(None, 64)	2112
dense_15 (Dense)	(None, 128)	8320
dense_16 (Dense)	(None, 784)	101136

Total params: 223,456

Trainable params: 223,456

Non-trainable params: 0

original



latent

representation



predicted

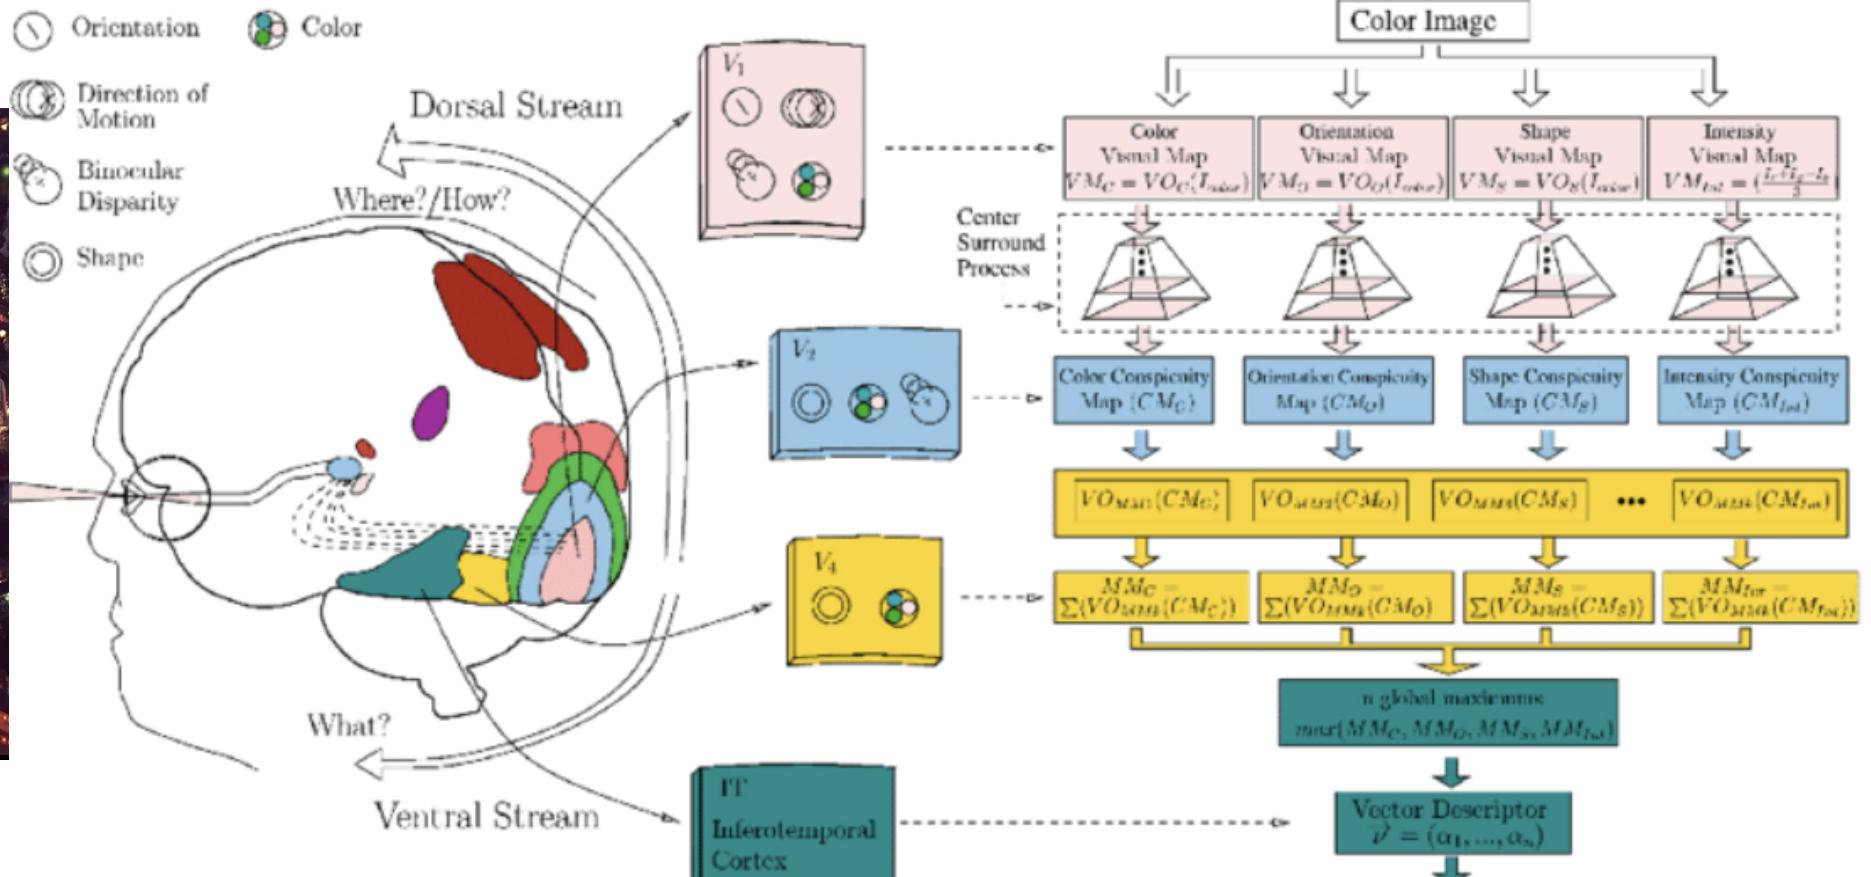




5/6
convolution

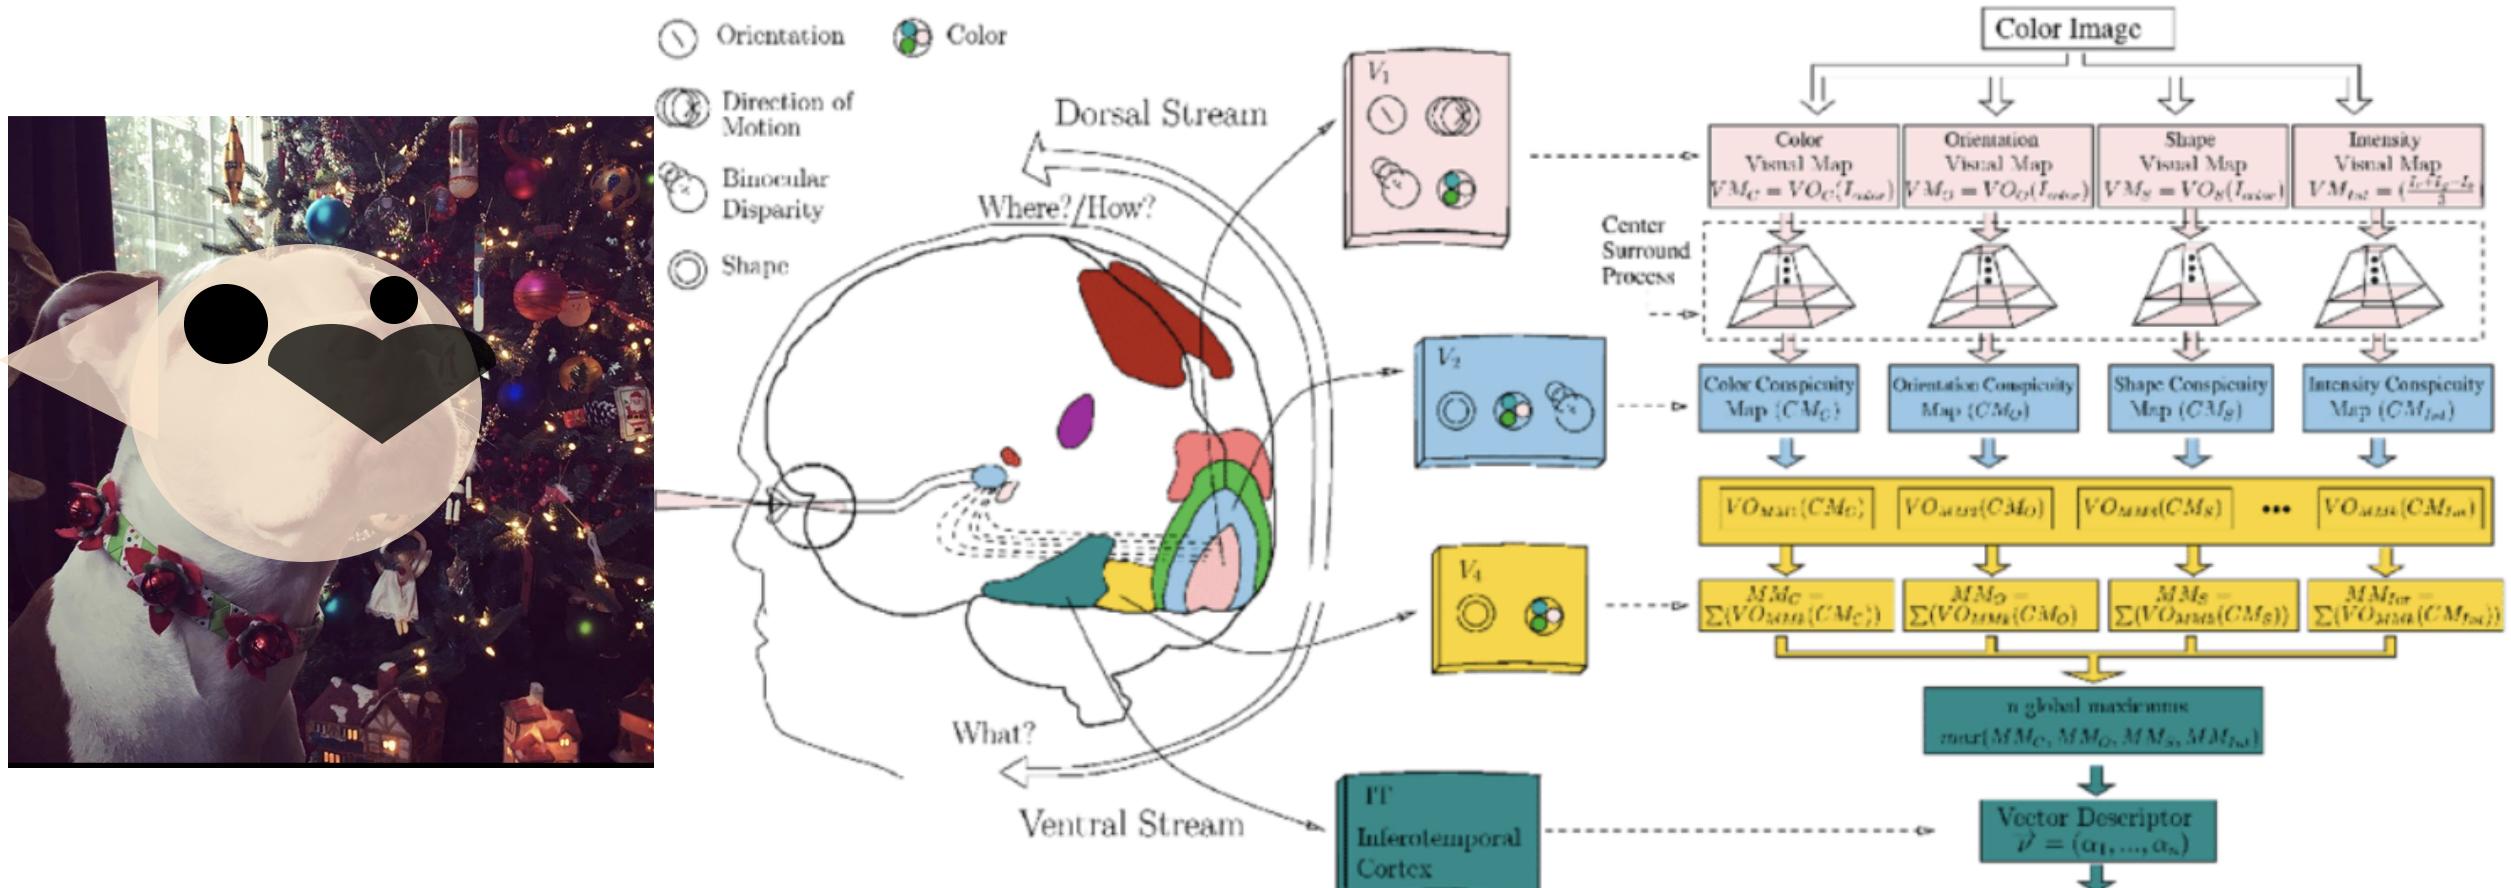


@akumadog



Brain Programming and the Random Search in Object Categorization
Oläglue et al 2017

The visual cortex learns hierarchically: first detects simple features, then more complex features and ensembles of features



CNN



CNN

1a

Convolution

Convolution

convolution is a mathematical operator on two functions

f and g

that produces a third function

$f \otimes g$

expressing *how the shape of one is modified by the other.*

Convolution Theorem

$$F(\nu) = \int_{\mathbb{R}^n} f(x) e^{-2\pi i x \cdot \nu} dx,$$

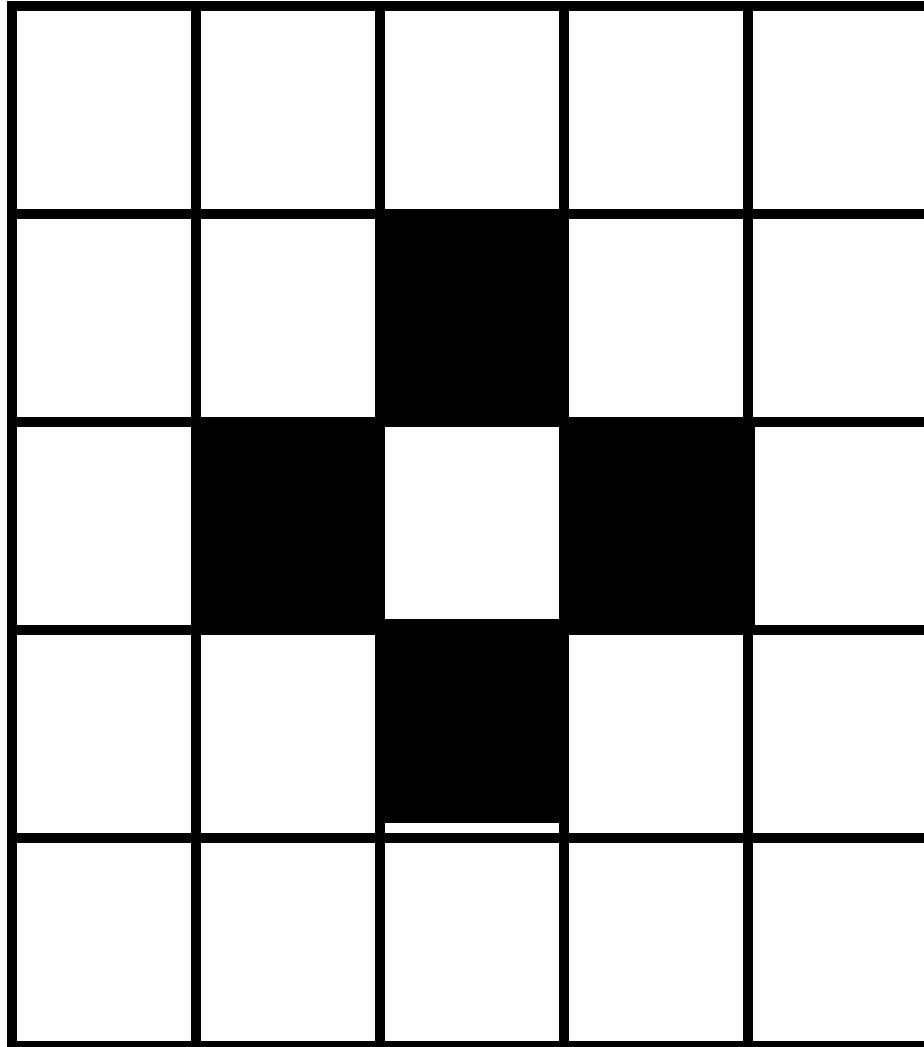
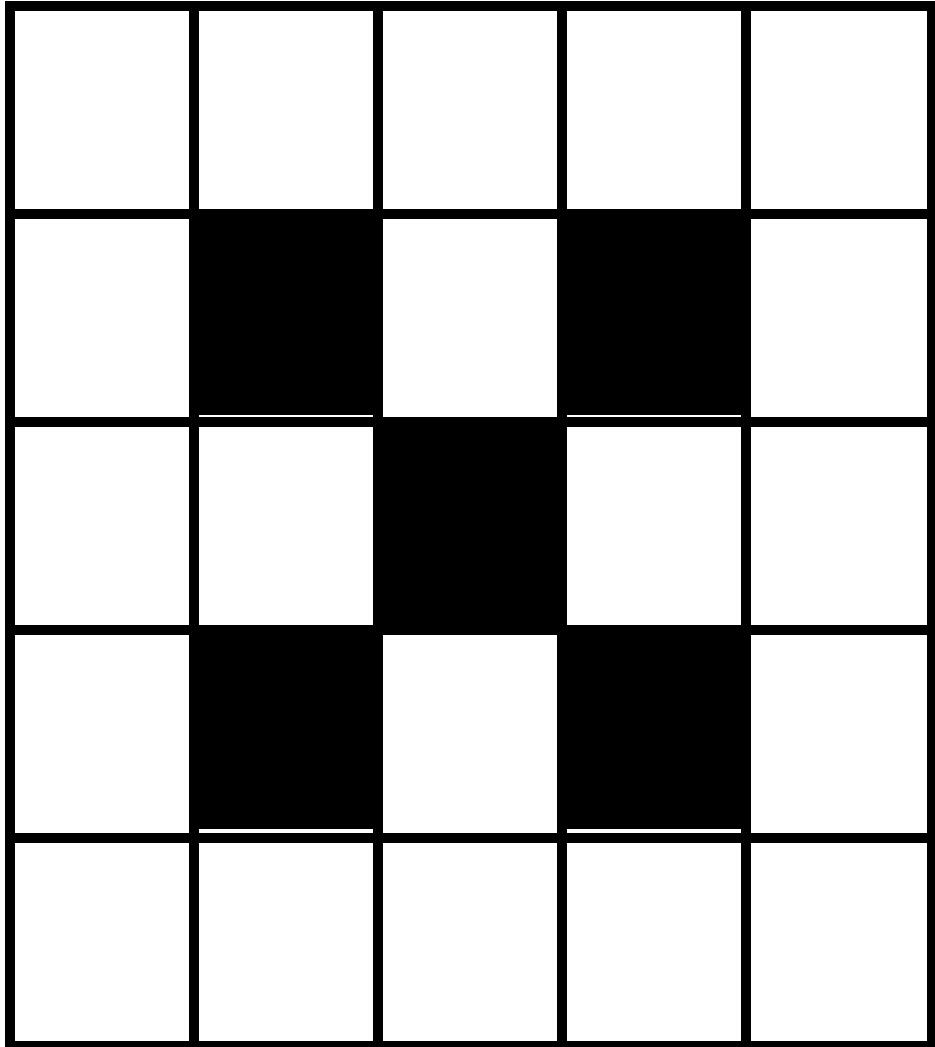
$$G(\nu) = \int_{\mathbb{R}^n} g(x) e^{-2\pi i x \cdot \nu} dx,$$

$$f * g = \mathcal{F}^{-1} \{ \mathcal{F}\{f\} \cdot \mathcal{F}\{g\} \}$$

\mathcal{F} fourier transform

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

two images.



```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1

-1	-1	-1	-1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1
-1	-1	-1	-1	-1

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1

feature maps

1	-1	-1
-1	1	-1
-1	-1	1

-1	-1	1
-1	1	-1
1	-1	-1

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1

 convolution

1	-1	-1
-1	1	-1
-1	-1	1

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

$$\begin{aligned} & (-1 * 1) + (-1 * -1) + (-1 * -1) + \\ & (-1 * -1) + (1 * 1) + (-1 * -1) \\ & (-1 * -1) + (-1 * -1) + (1 * 1) \\ & = 7 \end{aligned}$$

1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

7		

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

$$\begin{aligned} & (-1 * 1) + (-1 * -1) + (-1 * -1) + \\ & (-1 * 1) + (-1 * 1) + (-1 * 1) \\ & (-1 * -1) + (-1 * 1) + (-1 * 1) \\ & = -3 \end{aligned}$$

-1	-1	-1	-1	-1
-1	-1	-1	-1	-1
-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

7	-3	

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

7	-3	3

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
1	-1	-1	1	-1
-1	1	-1	-1	-1
-1	-1	1	1	-1
-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

7	-1	3
?		

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	-1	-1
-1	-1	1	-1	-1
-1	-1	-1	1	-1
-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

7	-1	3
?	?	

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	-1	-1
-1	-1	-1	-1	-1
-1	1	-1	-1	-1
-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

7	-1	3
?	?	

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	1	-1
1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

7	-1	3
?	?	

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1
-1	-1	-1	-1	-1
-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

7	-1	3
?	?	

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	1



1	-1	-1
-1	1	-1
-1	-1	1

=

7	-1	3
?	?	

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

input layer

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1

convolution layer



1	-1	-1
-1	1	-1
-1	-1	1

feature map

=

7	-1	3
-3		

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

input layer

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1

convolution layer



1	-1	-1
-1	1	-1
-1	-1	1

feature map

7	-3	3
-3	5	-3
3	-1	7

the feature map is "richer": we went from binary to R

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

input layer

-1	-1	-1	-1	-1
-1	1	-1	1	-1
-1	-1	1	-1	-1
-1	1	-1	1	-1
-1	-1	-1	-1	-1

convolution layer



1	-1	-1
-1	1	-1
-1	-1	1

feature map

=

7	-3	3
-3	5	-3
3	-1	7

the feature map is "richer": we went from
binary to R
and it is reminiscent of the original layer

Convolve with different feature: each neuron is 1 feature

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	-1
-1	1	-1
-1	-1	1

=

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



1	-1	1
-1	1	-1
1	-1	1

=

0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.11	0.33	-0.77	1.00	-0.77	0.33	-0.11
0.11	-0.55	0.55	-0.77	0.55	-0.55	0.11
-0.55	0.55	-0.55	0.33	-0.55	0.55	-0.55
0.33	-0.55	0.11	-0.11	0.11	-0.55	0.33

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1



-1	-1	1
-1	1	-1
1	-1	-1

=

0.33	-0.11	0.55	0.33	0.11	-0.11	0.77
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.77	-0.11	0.11	0.33	0.55	-0.11	0.33

CNN

tb

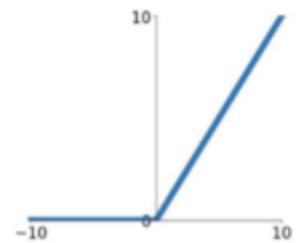
ReLU

```
1 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
```

ReLU: normalization that replaces negative values with 0's

7	-3	3
-3	5	-3
3	-1	7

ReLU
 $\max(0, x)$



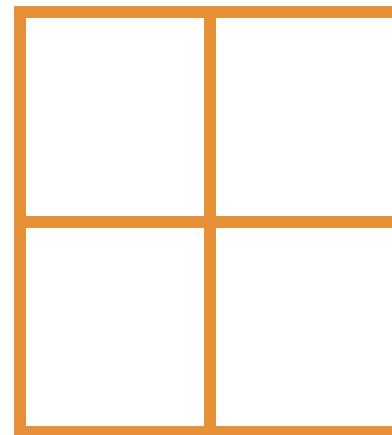
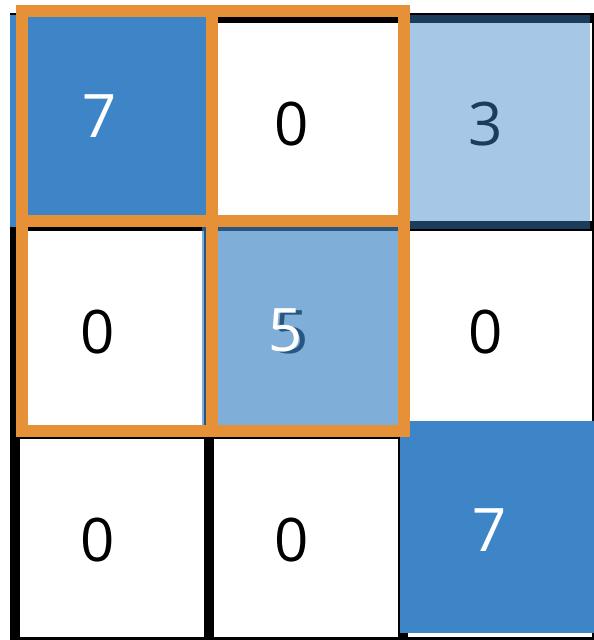
7	0	3
0	5	0
3	0	7

CNN

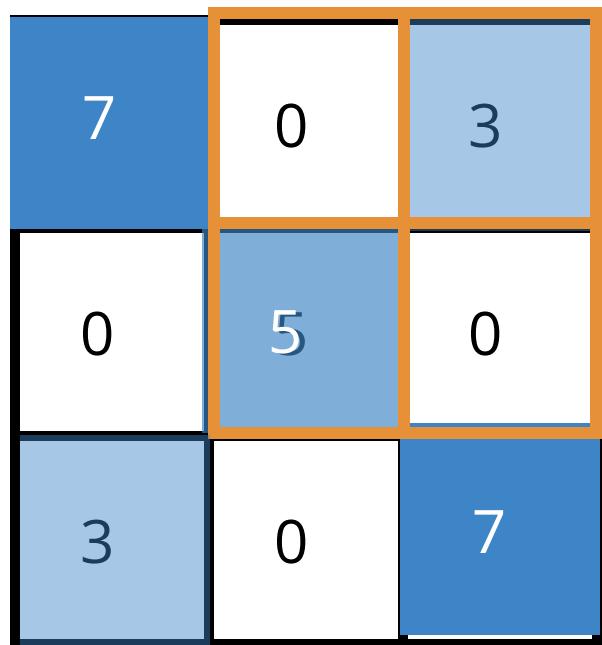
tc

Max-Pool

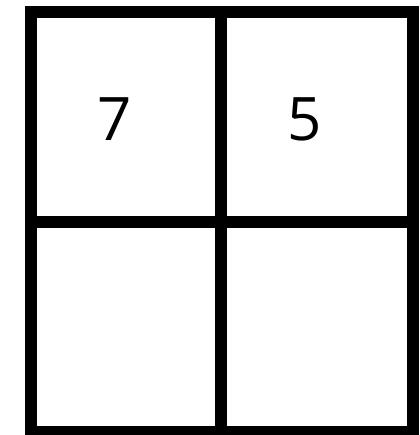
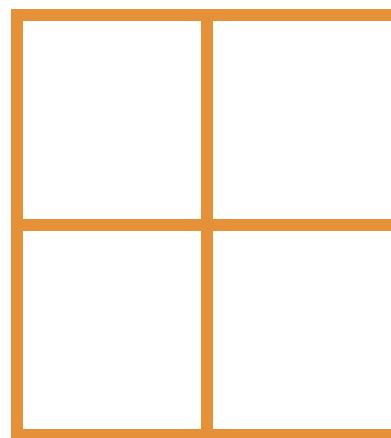
MaxPooling: reduce image size, generalizes result



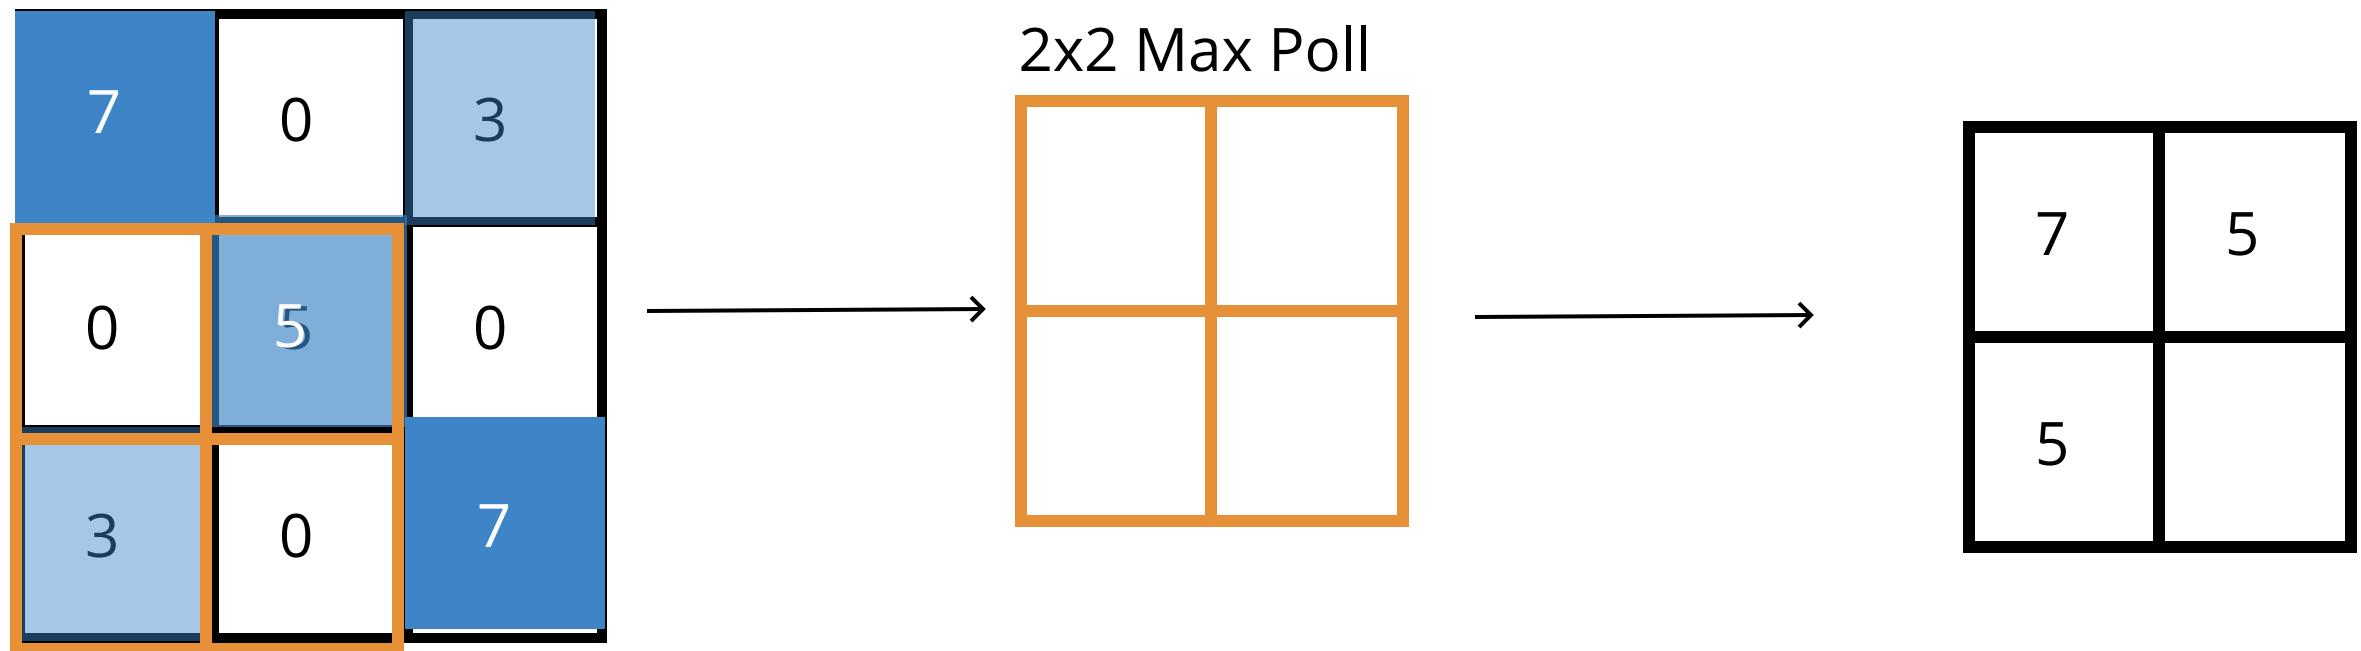
MaxPooling: reduce image size, generalizes result



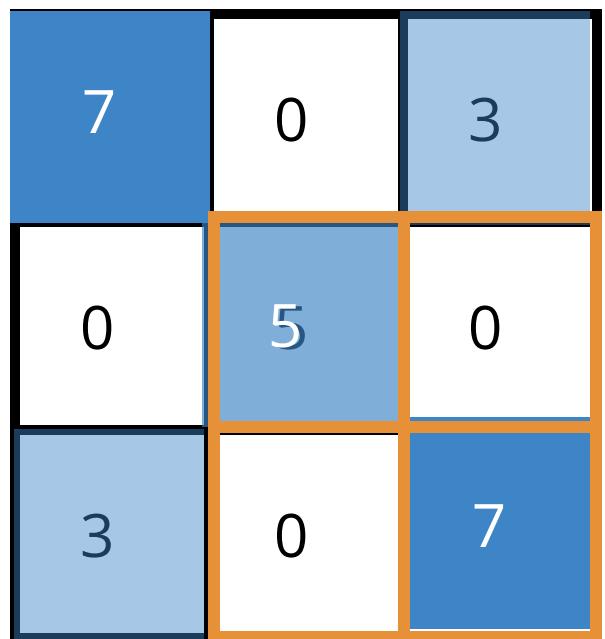
2x2 Max Poll



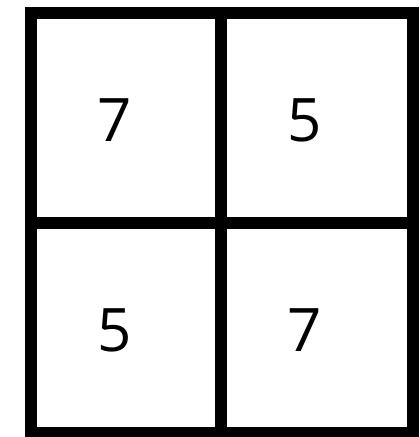
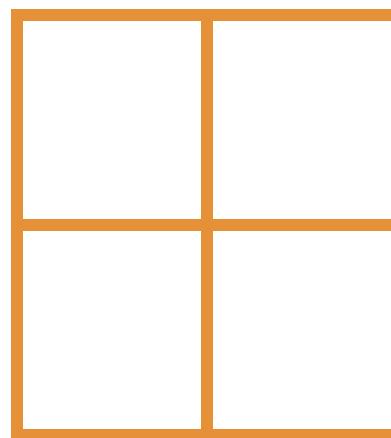
MaxPooling: reduce image size, generalizes result



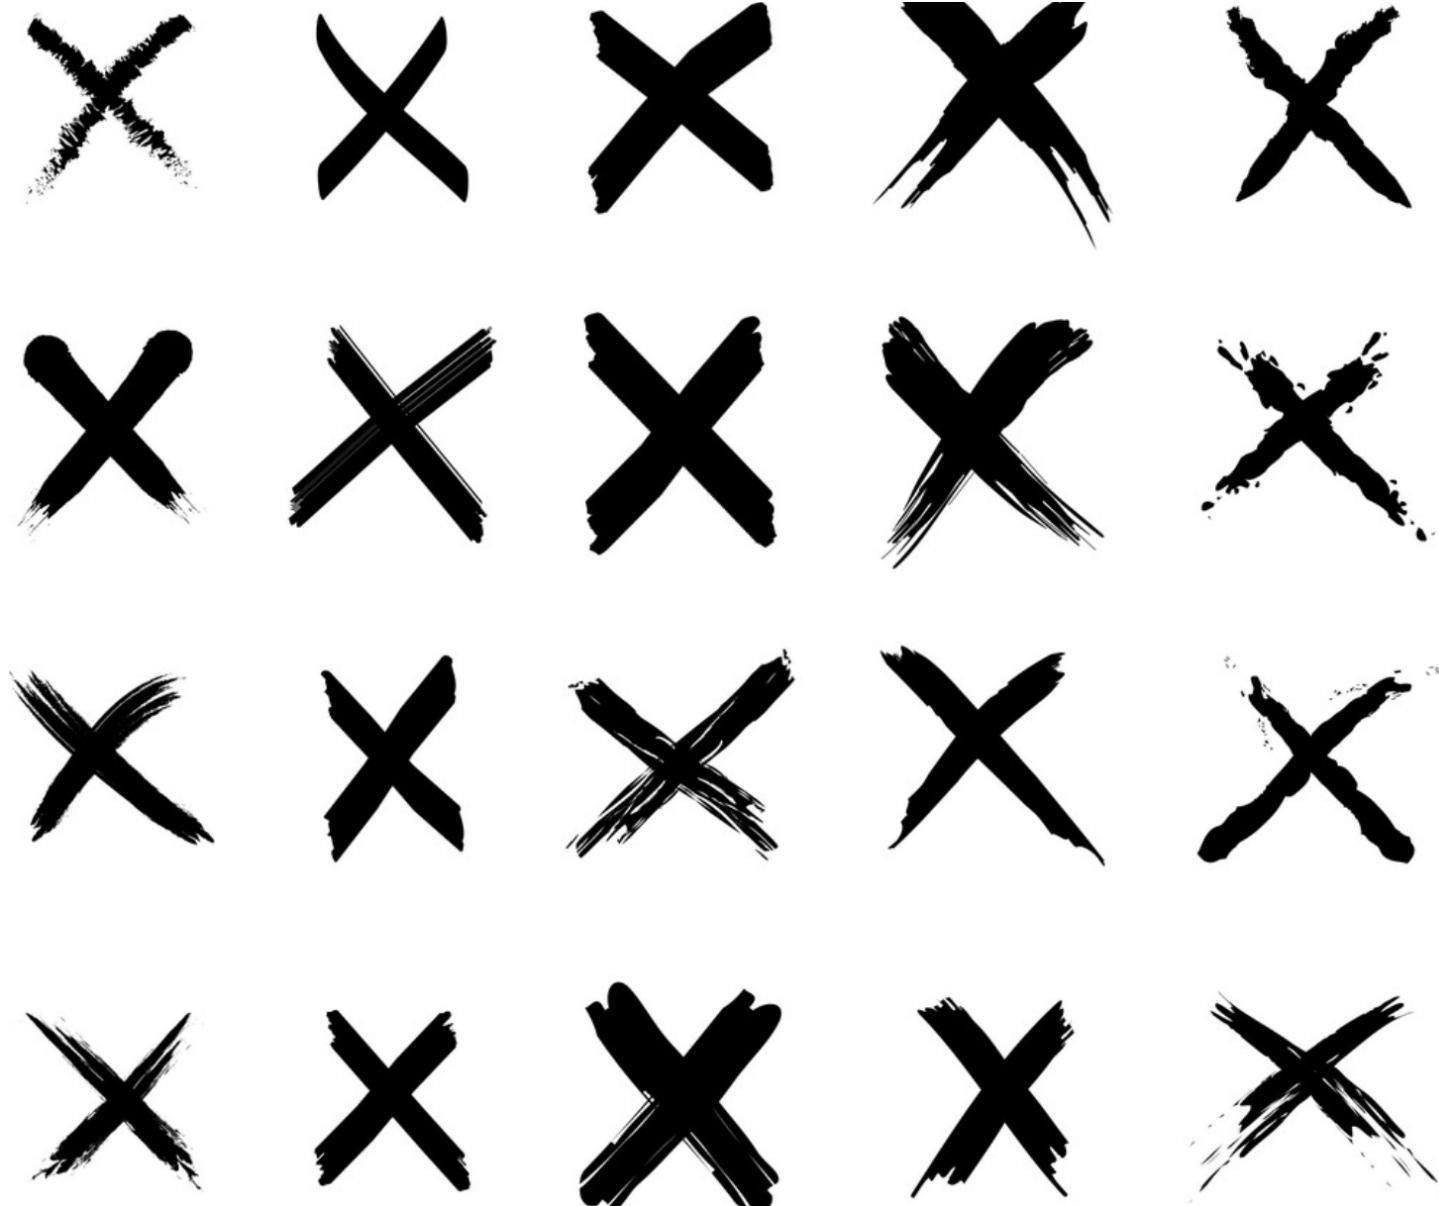
MaxPooling: reduce image size, generalizes result



2x2 Max Poll



```
1 model.add(MaxPooling2D(pool_size=(2, 2)))
```



MaxPooling:
reduce image
size & *generalizes*
result

By reducing the size and picking the maximum of a sub-region we make the network less sensitive to specific details

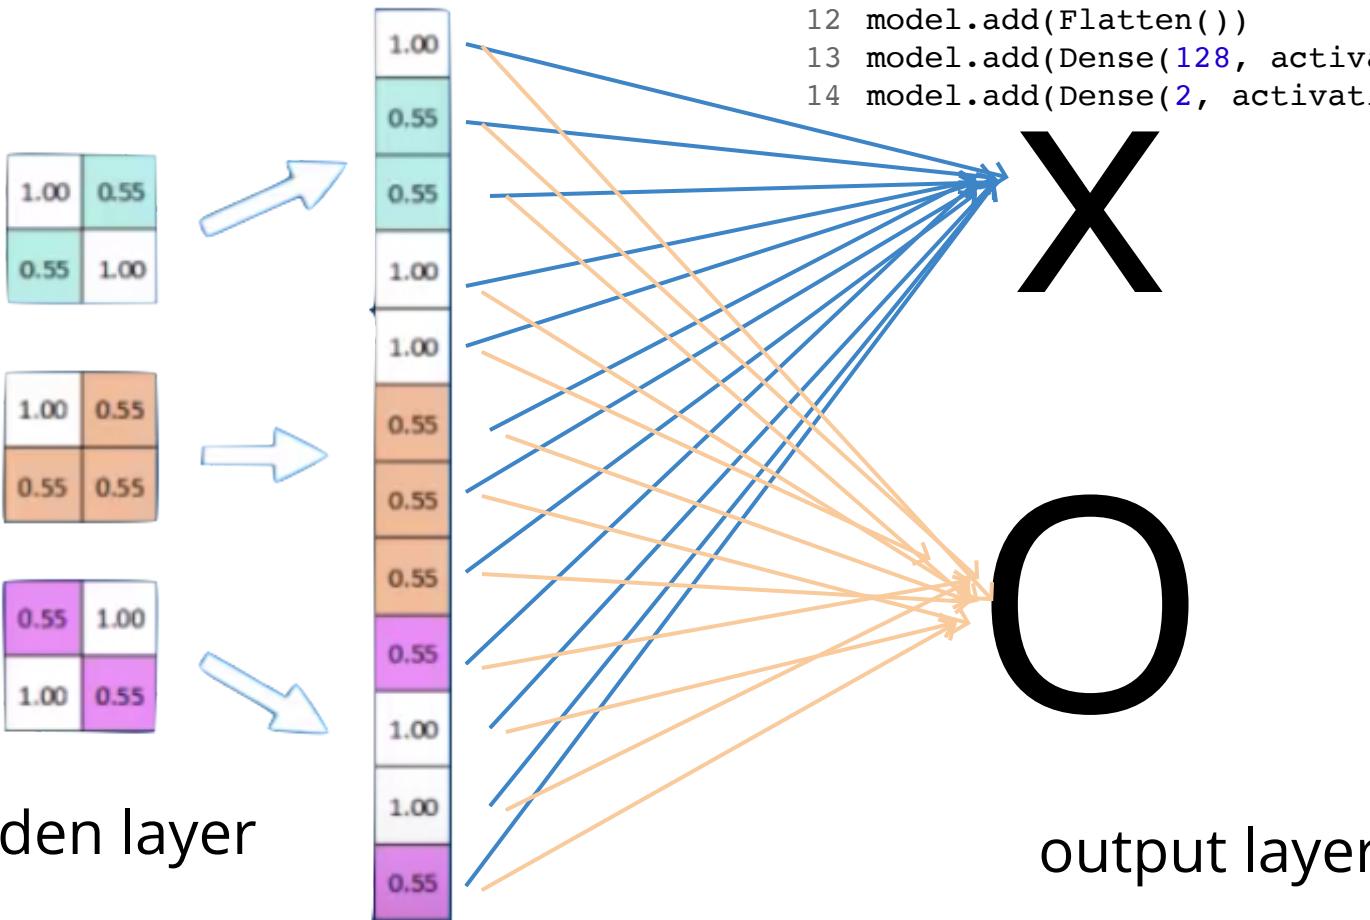
CNN

```
1 from keras.models import Sequential  
2 from keras.layers import Dense, Dropout, Flatten  
3 from keras.layers import Conv2D, MaxPooling2D  
4 model = Sequential()  
5 model.add(Conv2D(32, kernel_size=(10, 10),  
6                  activation='relu',  
7                  input_shape=input_shape))  
8 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))  
9 model.add(MaxPooling2D(pool_size=(2, 2)))  
10 model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))  
11 model.add(MaxPooling2D(pool_size=(2, 2)))
```

-1	-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1	-1

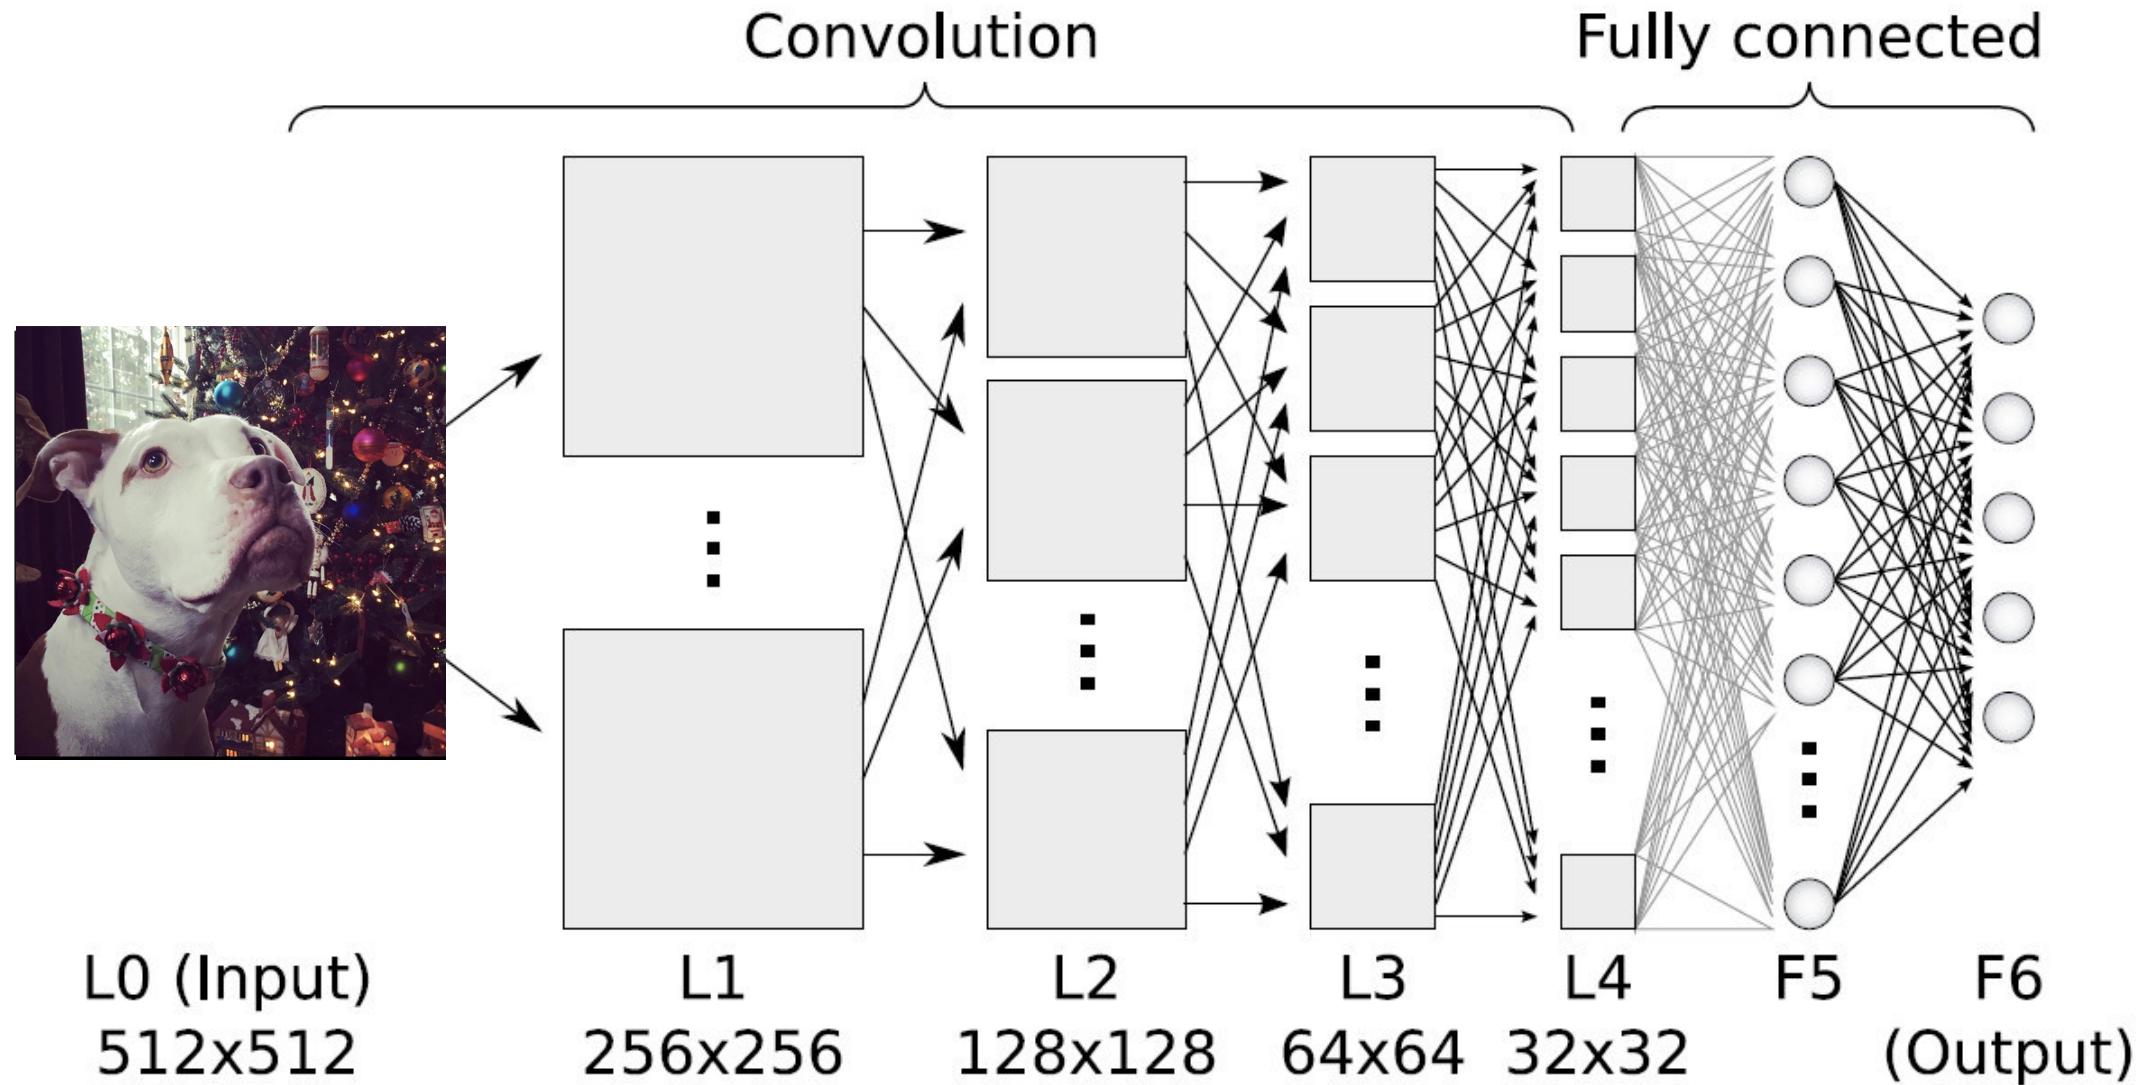


final layer: the final layer is fully connected MPL



```
1 from keras.models import Sequential
2 from keras.layers import Dense, Dropout, Flatten
3 from keras.layers import Conv2D, MaxPooling2D
4 model = Sequential()
5 model.add(Conv2D(32, kernel_size=(10, 10),
6                  activation='relu',
7                  input_shape=input_shape))
8 model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
9 model.add(MaxPooling2D(pool_size=(2, 2)))
10 model.add(Conv2D(32, kernel_size=(3, 3), activation='relu'))
11 model.add(MaxPooling2D(pool_size=(2, 2)))
12 model.add(Flatten())
13 model.add(Dense(128, activation='relu'))
14 model.add(Dense(2, activation='softmax'))
```

Stack multiple convolution layers



6/6

*Generative AI in
society*

Faces95 - Computer Vision Science Research Projects, Dr Libor Spacek.

2) **Faces95:** This dataset consists of 72 (male and female) individuals with 20 images per individual. The background of images consists of a red curtain and also has the background variation that caused by shadows. There are some expression and lighting variation.



Examples Faces95 images

White Males

White Female

Non-White Males

Non-White Female

Faces95 - Computer Vision Science Research Projects, Dr Libor Spacek.

2) **Faces95:** This dataset consists of 72 (male and female) individuals with 20 images per individual. The background of images consists of a red curtain and also has the background variation that caused by shadows. There are some expression and lighting variation.



Examples Faces95 images

White Males: 43

White Female: 6

Non-White Males: 12

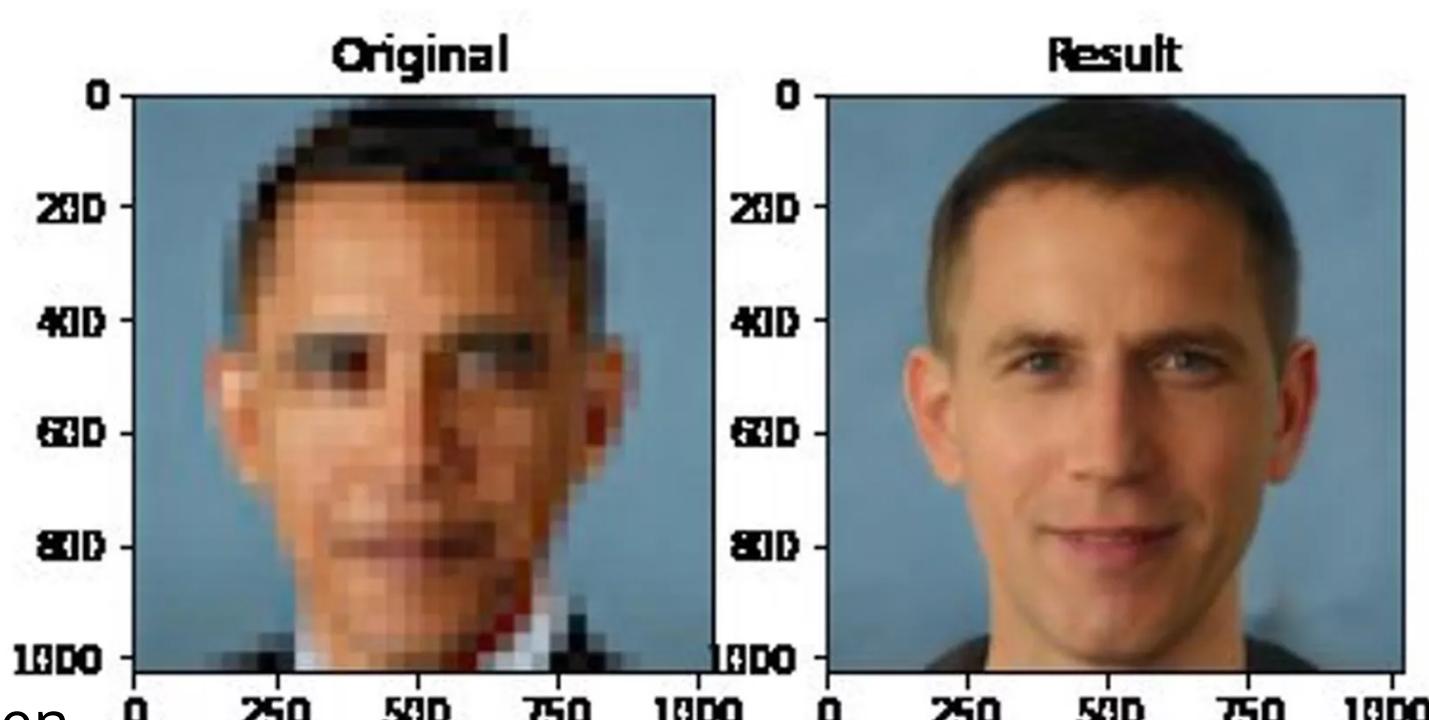
Non-White Female: 1

models are neutral, the bias is in the data (or is it?)

Why does this AI model whitens Obama face?

Simple answer: the data is biased. The algorithm is fed more images of white people

But really, would the opposite have been acceptable? *The bias is in society*



<https://www.theverge.com/21298762/face-depixelizer-ai-machine-learning-tool-pulse-stylegan-obama-bias>

models are neutral, the bias is in the data (or is it?)

Why does this AI model whitens Obama face?

Simple answer: the data is biased. The algorithm is fed more images of white people



<https://www.theverge.com/21298762/face-depixelizer-ai-machine-learning-tool-pulse-stylegan-obama-bias>

~~models are neutral, the bias is in the data~~

The bias is in the **data**

The bias is in the **models** and
the decision we make

The bias is in **how we choose to**
optimize our model

**The bias is society that
provides the framework to
validate our biased models**

Should AI reflect
who we are
(and *enforce and grow our bias*)
or should it reflect who we
aspire to be?
(*and who decides what that is?*)

none of this is new

~~models are neutral,~~
~~the bias is in the data~~

The bias is in the **data**

The bias is in the **models** and
the decision we make

The bias is in **how we choose to**
optimize our model

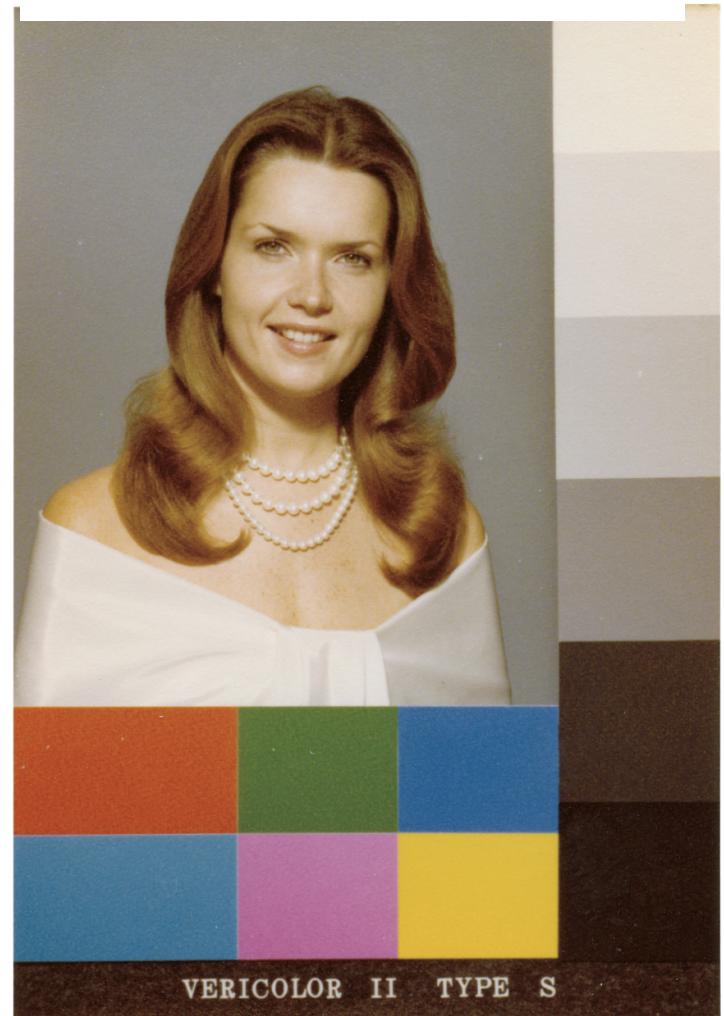
**The bias is society that
provides the framework to
validate our biased models**

Lens

LENS

The Racial Bias Built Into
Photography

Sarah Lewis explores the relationship between racism and the camera.



Shirley Card, 1978. Courtesy of Hermann Zschiegner

[https://www.nytimes.com/2019/04/25/lens/sar ah-lewis-racial-bias-photography.html](https://www.nytimes.com/2019/04/25/lens/sarah-lewis-racial-bias-photography.html)

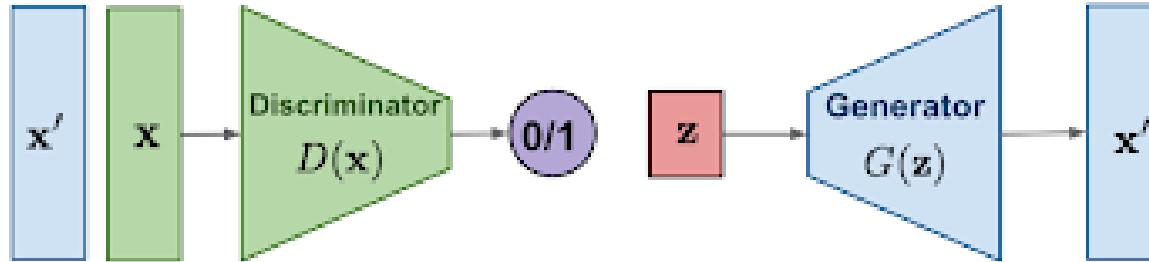
7/16

*In defense of
autoencoders*

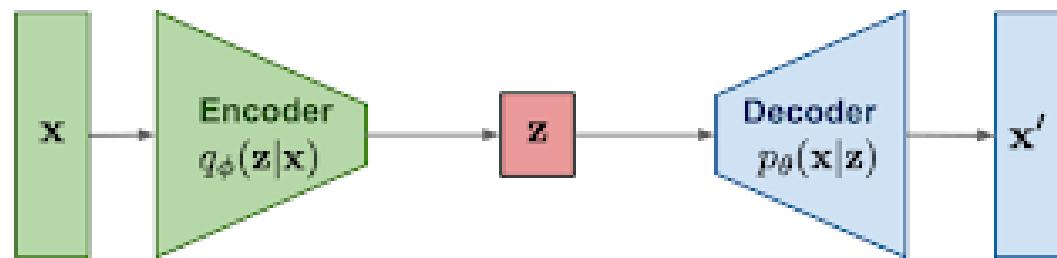
see also

<https://arxiv.org/pdf/2103.04922.pdf>

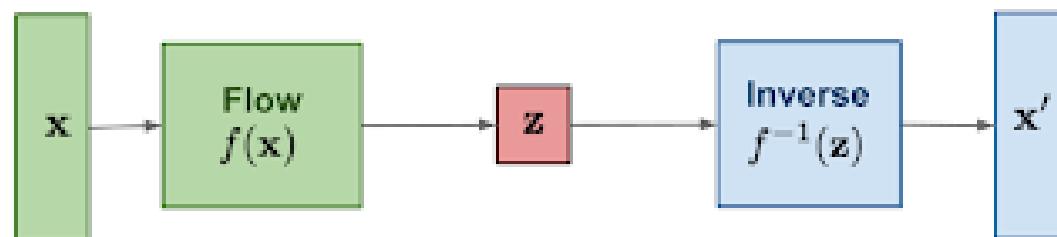
GAN: Adversarial training



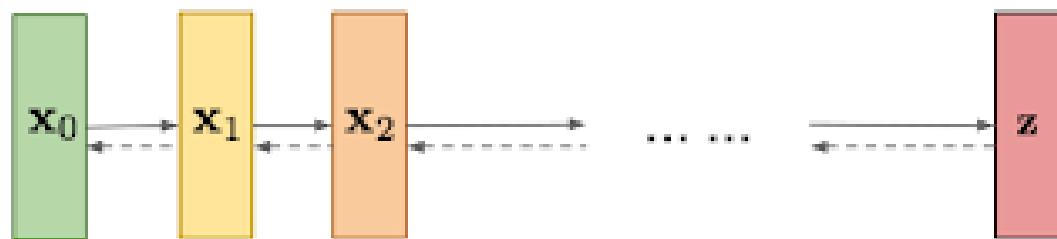
VAE: maximize variational lower bound



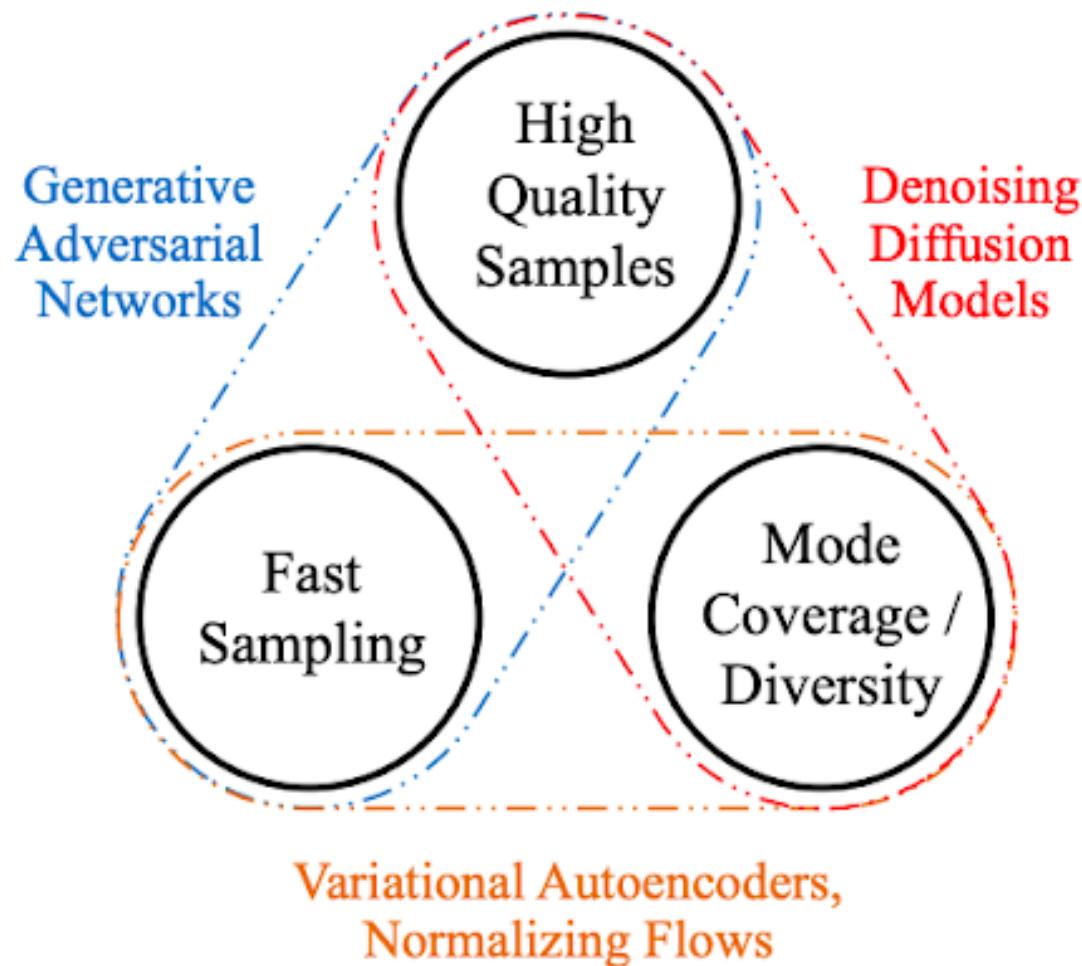
Flow-based models:
Invertible transform of distributions

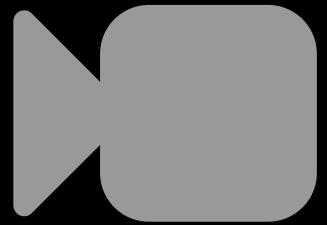


Diffusion models:
Gradually add Gaussian noise and then reverse



Which generative AI is right for you??





thank you!

federica bianco

University of Delaware

Department of Physics and
Astronomy

Biden School of Public

Policy and Administration
Data Science Institute



@fedhere

fbianco@udel.edu